



The ACM Conference Series on  
**Recommender Systems**

# The Recommender Problem *Revisited*

*Xavier Amatriain*  
Research/Engineering Director @  
Netflix



Xavier Amatriain – October 2014 – Recsys



# Index

1. The Recommender Problem
2. Traditional Recommendation Methods
  - 2.1. Collaborative Filtering
  - 2.2. Content-based Recommendations
  - 2.3. Hybrid Approaches
3. Beyond Traditional Methods
  - 3.1. Learning to Rank
  - 3.2. Similarity
  - 3.3. Deep Learning
  - 3.4. Social Recommendations
  - 3.5. Page Optimization
  - 3.6. Tensor Factorization and Factorization Machines
  - 3.7. MAB Explore/Exploit
4. References

# 1. The Recommender Problem



# Everything is a recommendation

The collage consists of several screenshots from different websites and interfaces:

- Huffington Post Article:** A news article titled "Netflix's New 'My List' Feature Knows You Better Than You Know Yourself (Because Algorithms)". The article is by Dino Grandoni and posted on August 21, 2013. It features a large image of the Netflix logo on a building.
- Facebook Photos:** A screenshot of a Facebook photo album titled "Our Trip to Yellowstone". It shows a person reaching out towards a waterfall. The caption reads: "We decided to go to Yellowstone for the weekend to reconnect with nature. It was a memorable trip with good friends."
- Facebook News Feed:** A sidebar showing a news feed with various posts and photos.
- Amazon Recommendations:** A sidebar showing recommended products, including "JULIA ALAN WHOLE ALLEN ANGELICA HUSTON MARK CRISTIN MANHATTAN MYSTERY".
- Twitter Header:** A screenshot of a Twitter header featuring a woman's face.
- LinkedIn Profile:** A screenshot of a LinkedIn profile page.
- Reddit Subreddit:** A screenshot of a Reddit subreddit page.
- YouTube Channel:** A screenshot of a YouTube channel page.
- Netflix Home Page:** A screenshot of the Netflix home page showing movie and TV show thumbnails.

**NETFLIX**

# The “Recommender problem”

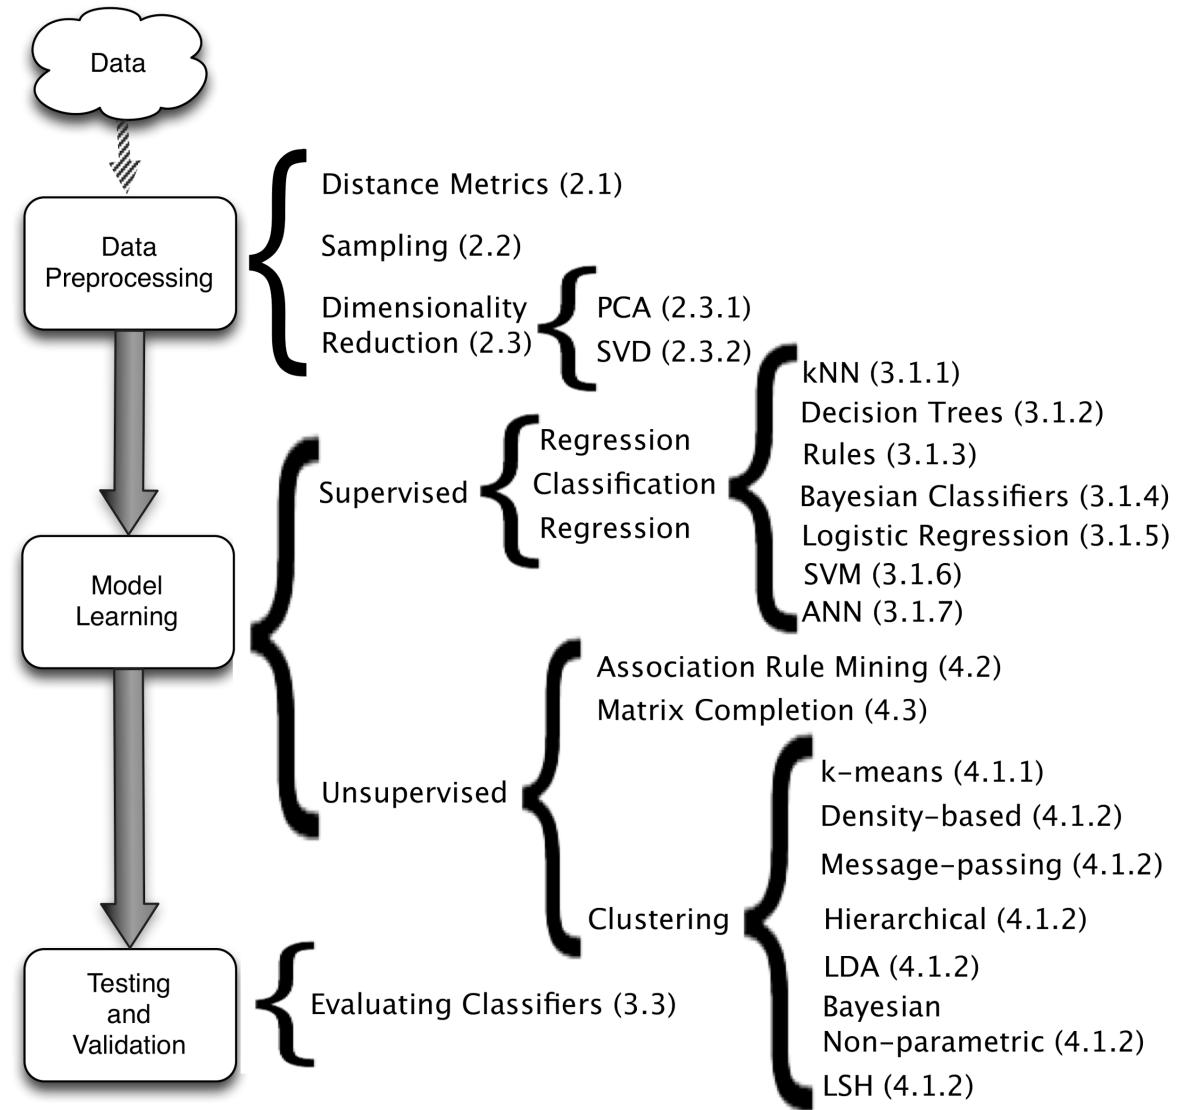
- *Traditional* definition: Estimate a utility function that automatically predicts how a user will like an item.
- Based on:
  - Past behavior
  - Relations to other users
  - Item similarity
  - Context
  - ...



# Recommendation as data mining

The core of the  
Recommendation  
Engine can be  
assimilated to a general  
data mining problem

(Amatriain et al. *Data Mining Methods for Recommender Systems in Recommender Systems Handbook*)



# Machine Learning + all those other things

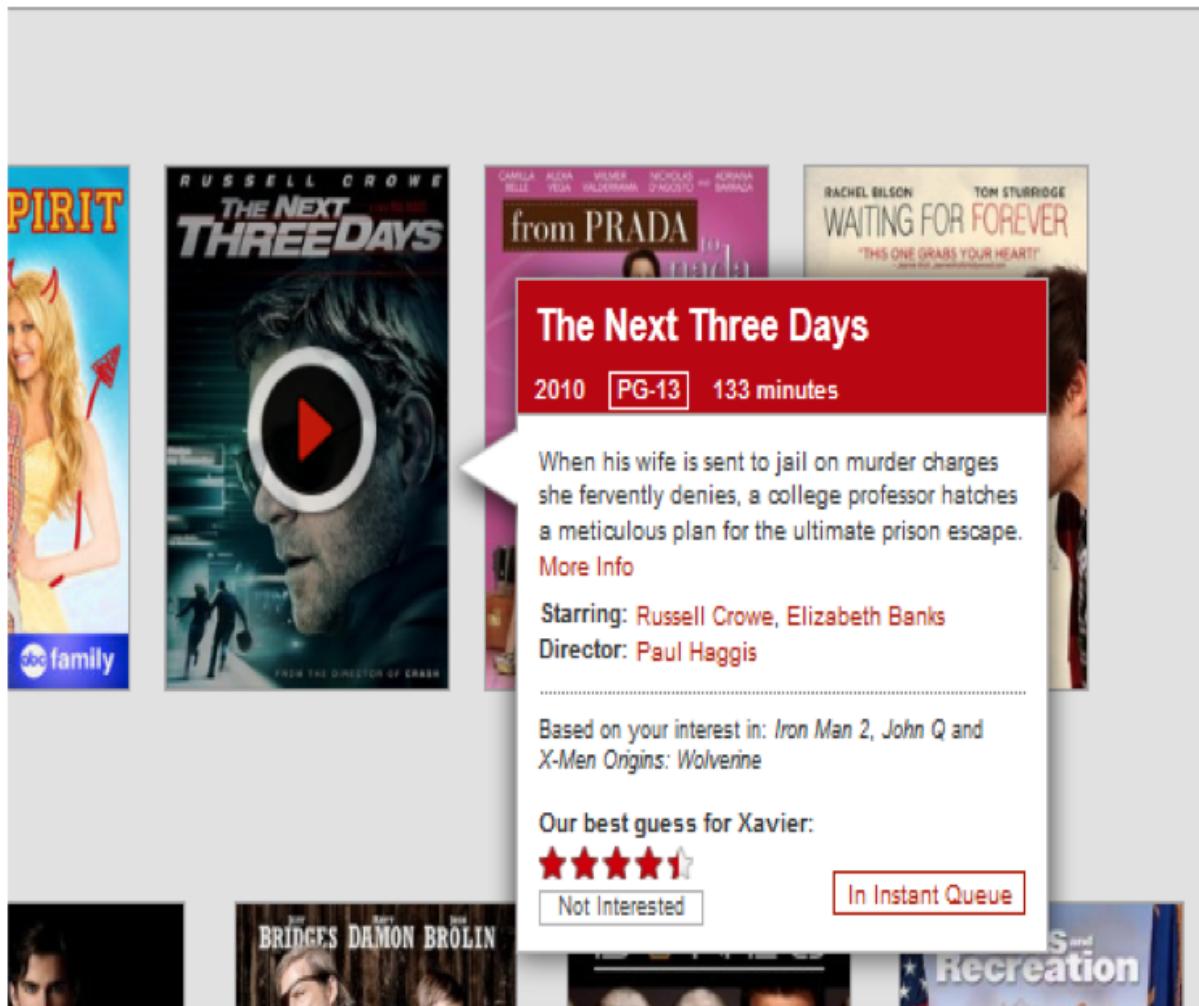
- User Interface
- System requirements (efficiency, scalability, privacy....)
- Serendipity
- Diversity
- Awareness
- Explanations
- ...



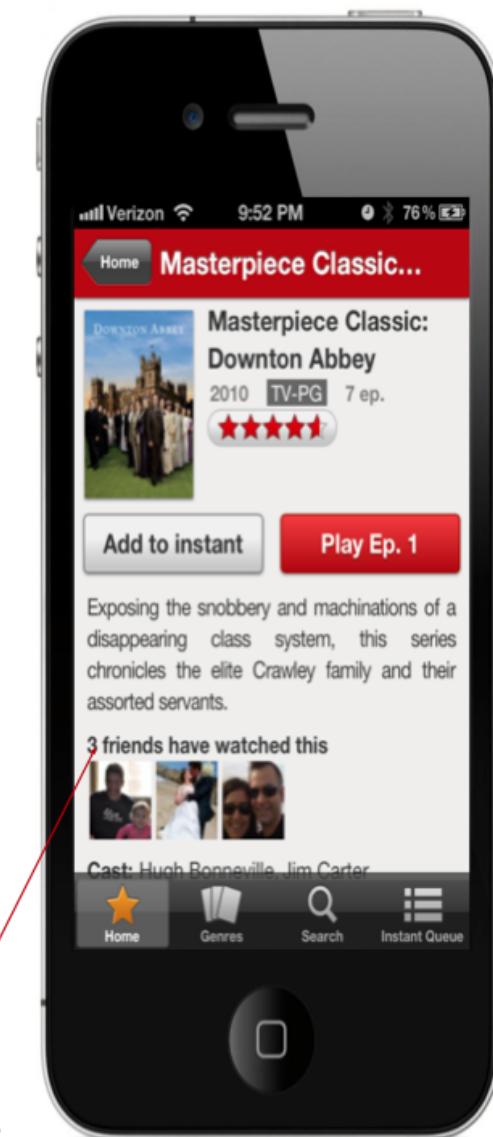
# Serendipity

- Unsought finding
- Don't recommend items the user already knows or **would have found anyway.**
- Expand the user's taste into neighboring areas by improving the obvious
- Collaborative filtering can offer controllable serendipity (e.g. controlling how many neighbors to use in the recommendation)

# Explanation/Support for Recommendations



Social Support



# Diversity & Awareness

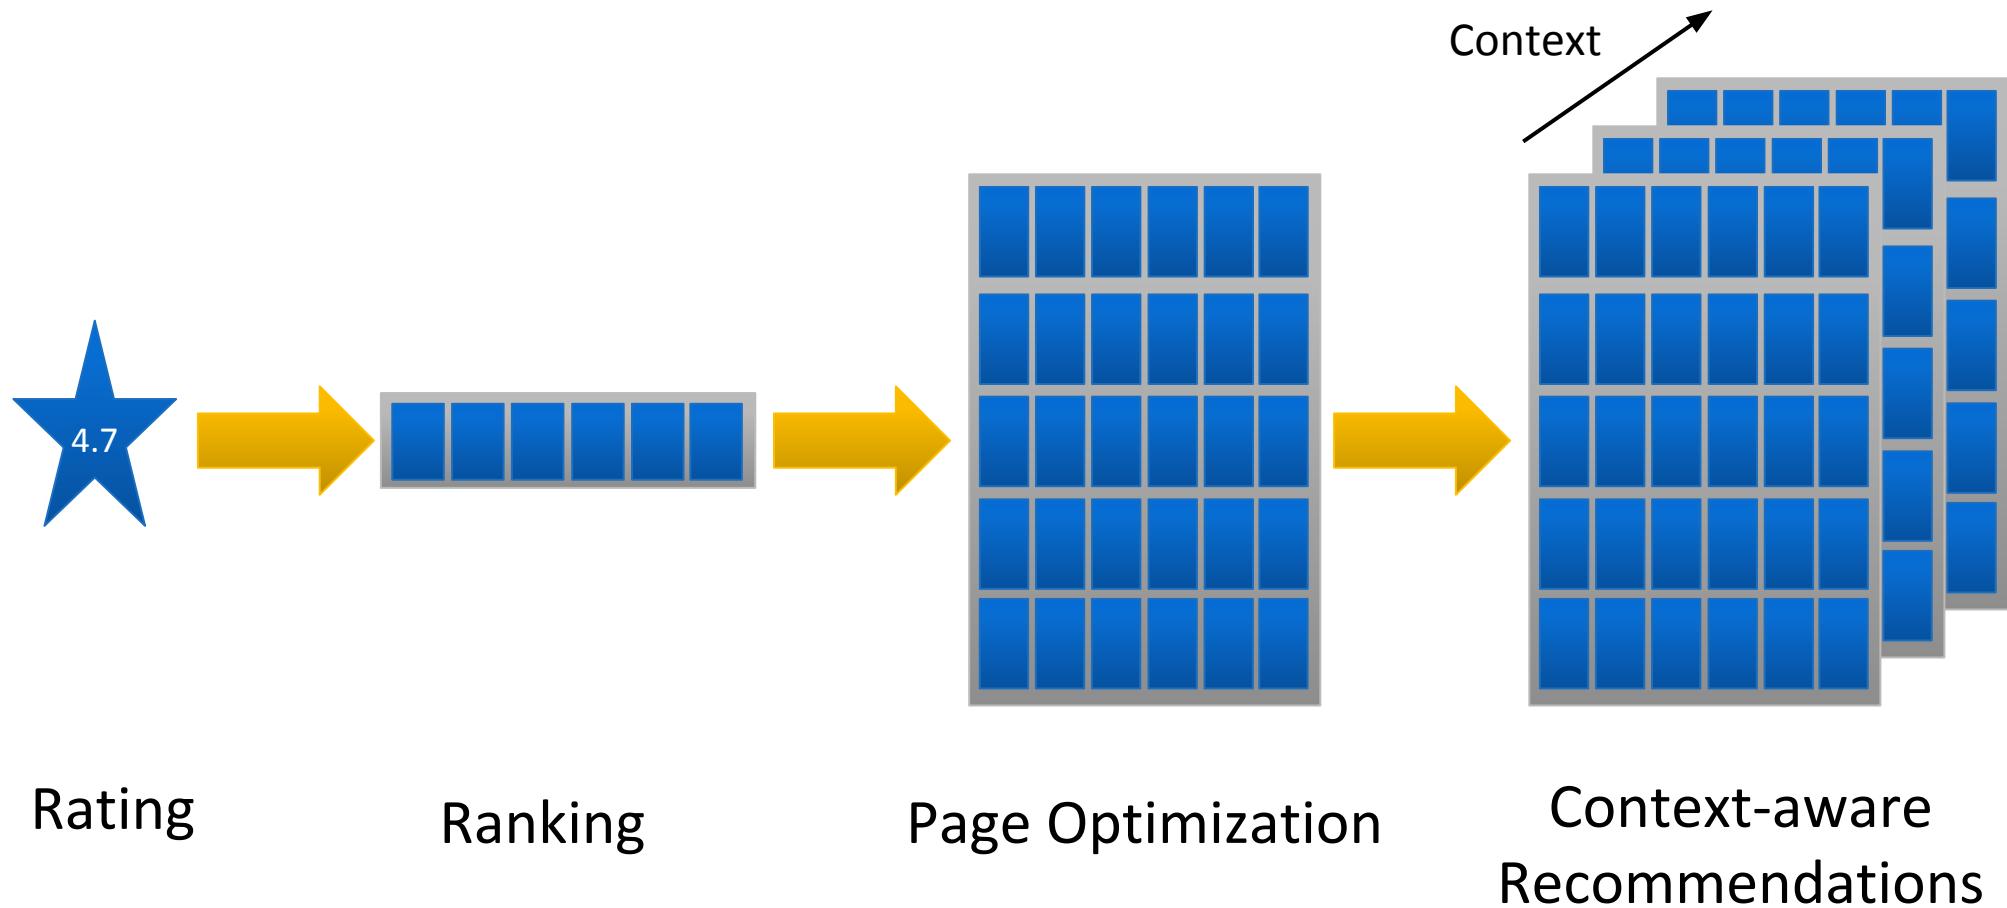
Personalization awareness



Diversity



# Evolution of the Recommender Problem



NETFLIX

# Index

1. The Recommender Problem
2. “Traditional” Methods
  - 2.1. Collaborative Filtering
  - 2.2. Content-based Recommendations
  - 2.3. Hybrid Approaches
3. Beyond Traditional Methods
  - 3.1. Learning to Rank
  - 3.2. Similarity
  - 3.3. Deep Learning
  - 3.4. Social Recommendations
  - 3.5. Page Optimization
  - 3.6. Tensor Factorization and Factorization Machines
  - 3.7. MAB Explore/Exploit
4. References

## 2. Traditional Approaches



## 2.1. Collaborative Filtering



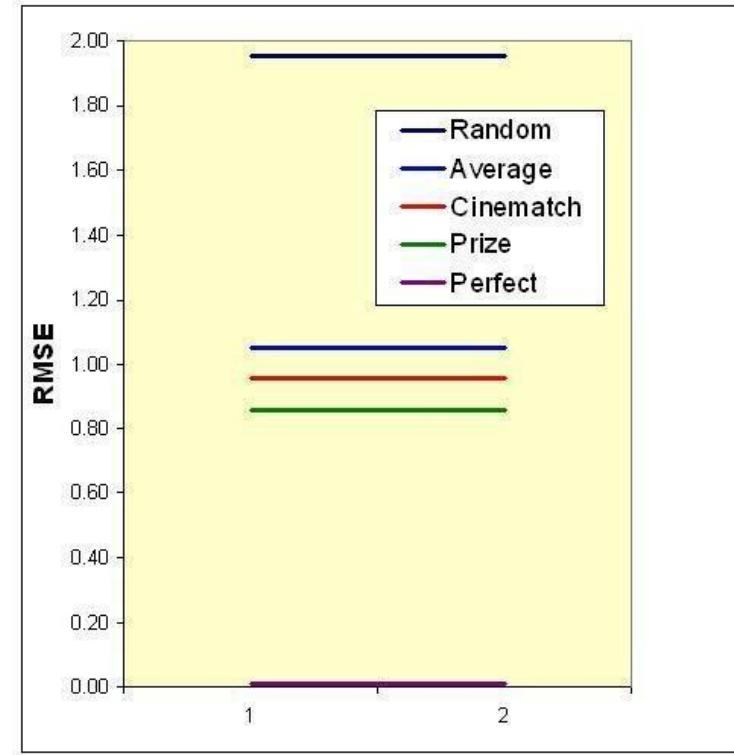
# Personalized vs Non-Personalised CF

- CF recommendations are **personalized**: prediction based only on **similar users**
- **Non-personalized** collaborative-based recommendation: average the recommendations of **ALL** the users
- How would the two approaches compare?



# Personalized vs. Not Personalized

- Netflix Prize: it is very simple to produce “reasonable” recommendations and extremely difficult to improve them to become “great”
- But there is a huge difference in business value between reasonable and great



Data Set	users	items	total	density	MAE Non Pers	MAE Pers
Jester	48483	100	3519449	0,725	0,220	0,152
MovieLens	6040	3952	1000209	0,041	0,233	0,179
EachMovie	74424	1649	2811718	0,022	0,223	0,151



# User-based CF

The basic steps:

1. Identify set of ratings for the **target/active user**
2. Identify set of users most similar to the target/active user according to a similarity function (**neighborhood** formation)
3. Identify the products these similar users liked
4. **Generate a prediction** - rating that would be given by the target user to the product - for each one of these products
5. Based on this predicted rating recommend a set of top N products



# Item-Based CF

- Look into the items the target user has rated
- Compute how **similar** they are to the target item
  - Similarity **only using** past **ratings** from other users
- Select k most similar items.
- Compute Prediction by taking weighted average on the target user's ratings on the most similar items.



# CF: Pros/Cons

- Requires **minimal knowledge**
- Produces good-enough results in most cases

Challenges:

- **Sparsity** – evaluation of large itemsets where user/item interactions are under 1%.
- **Scalability** - Nearest neighbor require computation that grows with both the number of users and the number of items.
- ...



# Model-based Collaborative Filtering



# Model Based CF Algorithms

- Memory based
  - Use the entire user-item database to generate a prediction.
  - Usage of statistical techniques to find the neighbors – e.g. nearest-neighbor.
- Model-based
  - First develop a model of user
  - Type of model:
    - Probabilistic (e.g. Bayesian Network)
    - Clustering
    - Rule-based approaches (e.g. Association Rules)
    - Classification
    - Regression
    - LDA
    - ...



# Model-based CF: What we learned from the Netflix Prize



# Netflix Prize

COMPLETED

What we were interested in:

- High quality *recommendations*

Proxy question:

- Accuracy in predicted rating
- Improve by 10% = \$1million!



$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_{j=1}^n (y_j - \hat{y}_j)^2}$$

NETFLIX

# 2007 Progress Prize

- Top 2 algorithms
  - SVD - Prize RMSE: 0.8914
  - RBM - Prize RMSE: 0.8990
- Linear blend Prize RMSE: 0.88
- Currently in use as part of Netflix' rating prediction component
- Limitations
  - Designed for 100M ratings, we have 5B ratings
  - Not adaptable as users add ratings
  - Performance issues



# SVD/MF

$$X [n \times m] = U [n \times r] S [r \times r] (V [m \times r])^T$$

$$\begin{matrix} X \\ \left( \begin{array}{cccc} x_{11} & x_{12} & \dots & x_{1n} \\ x_{21} & x_{22} & \dots & \\ \vdots & \vdots & \ddots & \\ x_{m1} & & & x_{mn} \end{array} \right) \\ m \times n \end{matrix} = \begin{matrix} U \\ \left( \begin{array}{ccc} u_{11} & \dots & u_{1r} \\ \vdots & \ddots & \\ u_{m1} & & u_{mr} \end{array} \right) \\ m \times r \end{matrix} \begin{matrix} S \\ \left( \begin{array}{ccc} s_{11} & 0 & \dots \\ 0 & \ddots & \\ \vdots & & s_{rr} \end{array} \right) \\ r \times r \end{matrix} \begin{matrix} V^T \\ \left( \begin{array}{ccc} v_{11} & \dots & v_{1n} \\ \vdots & \ddots & \\ v_{r1} & & v_{rn} \end{array} \right) \\ r \times n \end{matrix}$$

- **X**:  $m \times n$  matrix (e.g., m users, n videos)
- **U**:  $m \times r$  matrix (m users, r factors)
- **S**:  $r \times r$  diagonal matrix (strength of each ‘factor’) (r: rank of the matrix)
- **V**:  $r \times n$  matrix (n videos, r factor)

# Simon Funk's SVD

- One of the most interesting findings during the Netflix Prize came out of a blog post
- Incremental, iterative, and approximate way to compute the SVD using gradient descent

Monday, December 11, 2006

*Netflix Update: Try This at Home*



not to be tied for third place on the [netflix prize](#). And I don't mean a sordid tale of computing in the jung  
the top ten or so.



# SVD for Rating Prediction

- User factor vectors  $p_u \in \Re^f$  and item-factors vector  $q_v \in \Re^f$
- Baseline (bias)  $b_{uv} = \mu + b_u + b_v$  (user & item deviation from average)
- Predict rating as  $r'_{uv} = b_{uv} + p_u^T q_v$
- **Asymmetric SVD** (Koren et. Al) asymmetric variation w. implicit feedback

$$r'_{uv} = b_{uv} + q_v^T \left( \left| R(u) \right|^{\frac{-1}{2}} \sum_{j \in R(u)} (r_{uj} - b_{uj}) x_j + \left| N(u) \right|^{\frac{-1}{2}} \sum_{j \in N(u)} y_j \right)$$

- Where
  - $q_v, x_v, y_v \in \Re^f$  are three item factor vectors
  - Users are not parametrized, but rather represented by:
    - $R(u)$ : items rated by user  $u$
    - $N(u)$ : items for which the user has given implicit preference (e.g. rated vs. not rated)



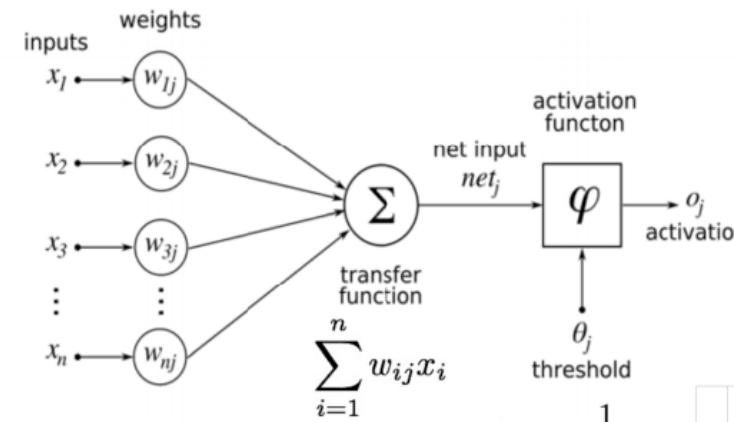
# Restricted Boltzmann Machines

- Each unit is a state that can be active or not active
- Each input to a unit is associated to a weight
- The transfer function  $\Sigma$  calculates a score for every unit based on the weighted sum of inputs
- Score is passed to the activation function  $\phi$  that calculates the probability of the unit to be active
- Restrict the connectivity to make learning easier.

Only one layer of hidden units.

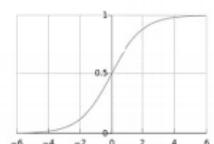
No connections between hidden units.

Hidden units are independent given visible states



$$\phi = \frac{1}{1 + e^{-\Sigma}}$$

$$P(o_j = 1|x) = \phi \left( \sum_{i=1}^n w_{ij}x_i + \theta_j \right)$$



# RBM for Recommendations

- Each visible unit = an item
- Num. of hidden units  $a$  is parameter
- In training phase, for each user:
  - If user rated item,  $v_i$  is activated
  - Activation states of  $v_i$  = inputs to  $h_j$
  - Based on activation,  $h_j$  is computed
  - Activation state of  $h_j$  becomes input to  $v_i$
  - Activation state of  $v_i$  is recalculated
  - Difference between current and past activation state for  $v_i$  used to update weights  $w_{ij}$  and thresholds
- In prediction phase:
  - For the items of the user the  $v_i$  are activated
  - Based on this the state of the  $h_j$  is computed
  - The activation of  $h_j$  is used as input to recompute the state of  $v_i$
  - Activation probabilities are used to recommend items

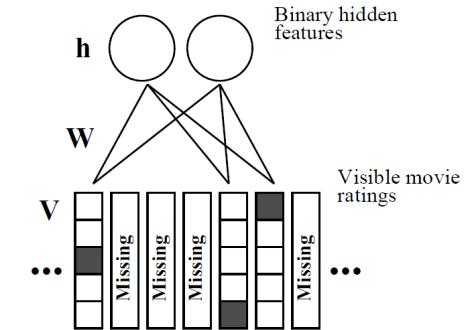
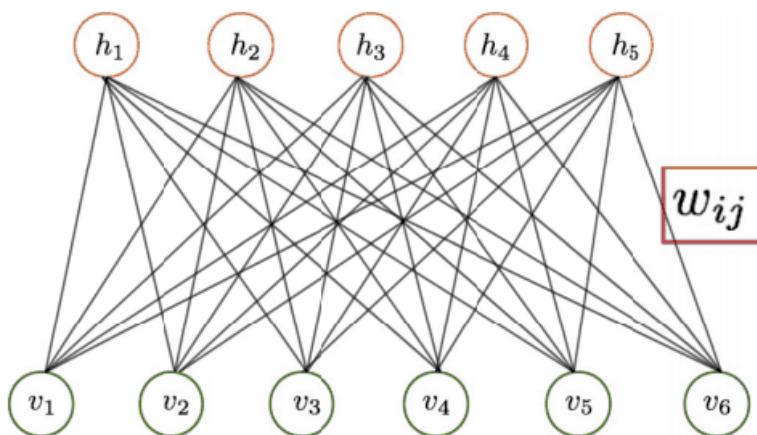


Figure 1. A restricted Boltzmann machine with binary hidden units and softmax visible units. For each user, the RBM only includes softmax units for the movies that user has rated. In addition to the symmetric weights between each hidden unit and each of the  $K = 5$  values of a softmax unit, there are 5 biases for each softmax unit and one for each hidden unit. When modeling user ratings with an RBM that has Gaussian hidden units, the top layer is composed of linear units with Gaussian noise.



# What about the final prize ensembles?

- Remember that current production model includes an **ensemble** of both SVD++ and RBMs
- Our offline studies showed final ensembles were too computationally intensive to scale
- Expected improvement not worth the engineering effort
- Plus.... Focus had already shifted to other issues that had more impact than rating prediction.

# Clustering



# Clustering

- Goal: **cluster** users and compute per-cluster “typical” preferences
- Users receive recommendations computed at the cluster level

# Locality-sensitive Hashing (LSH)

- Method for grouping similar items in highly dimensional spaces
- Find a hashing function s.t. similar items are grouped in the same buckets
- Main application is Nearest-neighbors
  - Hashing function is found iteratively by concatenating random hashing functions
  - Addresses one of NN main concerns: performance

# Other “interesting” clustering techniques

- k-means and all its variations
- Affinity Propagation
- Spectral Clustering
- LDA
- Non-parametric Bayesian Clustering (e.g. Chinese Restaurant Processes, HDPs)

# Association Rules



# Association rules

- Past purchases are interpreted as transactions of “associated” items
- If a visitor has some interest in Book 5, she will be recommended to buy Book 3 as well
- Recommendations are constrained to some minimum levels of confidence
- Fast to implement and execute (e.g. A Priori algorithm)

		Also bought ..					
		Book1	Book2	Book3	Book4	Book5	Book6
Customers who bought :	Book1				1	1	
	Book2			2		1	1
	Book3		2			2	
	Book4	1					
	Book5	1	1	2			
	Book6		1				



# Classifiers



# Classifiers

- **Classifiers** are general computational models trained using positive and negative examples
- They may take in inputs:
  - Vector of item features (action / adventure, Bruce Willis)
  - Preferences of customers (like action / adventure)
  - Relations among item
- E.g. Logistic Regression, Bayesian Networks, Support Vector Machines, Decision Trees, etc...

# Classifiers

- Classifiers can be used in CF and CB Recommenders
- Pros:
  - Versatile
  - Can be combined with other methods to improve accuracy of recommendations
- Cons:
  - Need a relevant training set
  - May overfit (Regularization)
- E.g. Logistic Regression, Bayesian Networks, Support Vector Machines, Decision Trees, etc...

# Limitations of Collaborative Filtering



# Limitations of Collaborative Filtering

- **Cold Start:** There needs to be enough other users already in the system to find a match. New items need to get enough ratings.
- **Popularity Bias:** Hard to recommend items to someone with unique tastes.
  - Tends to recommend popular items (items from the tail do not get so much data)



# Index

1. The Recommender Problem
2. “Traditional” Methods
  - 2.1. Collaborative Filtering
  - 2.2. Content-based Recommendations
  - 2.3. Hybrid Approaches
3. Beyond Novel Methods
  - 3.1. Learning to Rank
  - 3.2. Similarity
  - 3.3. Deep Learning
  - 3.4. Social Recommendations
  - 3.5. Page Optimization
  - 3.6. Tensor Factorization and Factorization Machines
  - 3.7. MAB Explore/Exploit
4. References

## 2.2 Content-based Recommenders



# Content-Based Recommendation

- Recommendations based on content of items rather than on other users' opinions/interactions
- Goal: recommend items **similar** to those the user liked
- Common for recommending **text-based products** (web pages, usenet news messages, )
- Items to recommend are “described” by their associated **features** (e.g. keywords)
- **User Model** structured in a “similar” way as the content: features/keywords more likely to occur in the preferred documents (lazy approach)
- The user model can be a **classifier** based on whatever technique (Neural Networks, Naïve Bayes...)



# Pros/cons of CB Approach

## Pros

- No need for data on other users: No cold-start or sparsity
- Able to recommend to users with unique tastes.
- Able to recommend new and unpopular items
- Can provide explanations by listing content-features

## Cons

- Requires content that can be encoded as meaningful features (difficult in some domains/catalogs)
- Users represented as learnable function of content features.
- Difficult to implement serendipity
- Easy to overfit (e.g. for a user with few data points)

# A word of caution

## Recommending New Movies: Even a Few Ratings Are More Valuable Than Metadata

István Pilászy \*

Dept. of Measurement and Information Systems  
Budapest University of Technology and  
Economics  
Magyar Tudósok krt. 2.  
Budapest, Hungary  
[pila@mit.bme.hu](mailto:pila@mit.bme.hu)

Domonkos Tikk \*,<sup>†</sup>

Dept. of Telecom. and Media Informatics  
Budapest University of Technology and  
Economics  
Magyar Tudósok krt. 2.  
Budapest, Hungary  
[tikk@tmit.bme.hu](mailto:tikk@tmit.bme.hu)

### ABSTRACT

The Netflix Prize (NP) competition gave much attention to collaborative filtering (CF) approaches. Matrix factorization (MF) based CF approaches assign low dimensional feature vectors to users and items. We link CF and content-based filtering (CBF) by finding a linear transformation that transforms user or item descriptions so that they are as close as possible to the feature vectors generated by MF for CF.

We propose methods for explicit feedback that are able to handle 140 000 features when feature vectors are very sparse. With movie metadata collected for the NP movies we show that the prediction performance of the methods is comparable to that of CF, and can be used to predict user preferences on new movies.

We also investigate the value of movie metadata compared to movie ratings in regards of predictive power. We compare

### 1. INTRODUCTION

The goal of recommender systems is to give personalized recommendation on items to users. Typically the recommendation is based on the former and current activity of the users, and metadata about users and items, if available.

There are two basic strategies that can be applied when generating recommendations. Collaborative filtering (CF) methods are based only on the activity of users, while content-based filtering (CBF) methods use only metadata. In this paper we propose hybrid methods, which try to benefit from both information sources.

The two most important families of CF methods are matrix factorization (MF) and neighbor-based approaches. Usually, the goal of MF is to find a low dimensional representation for both users and movies, i.e. each user and movie is associated with a feature vector. Movie metadata (which



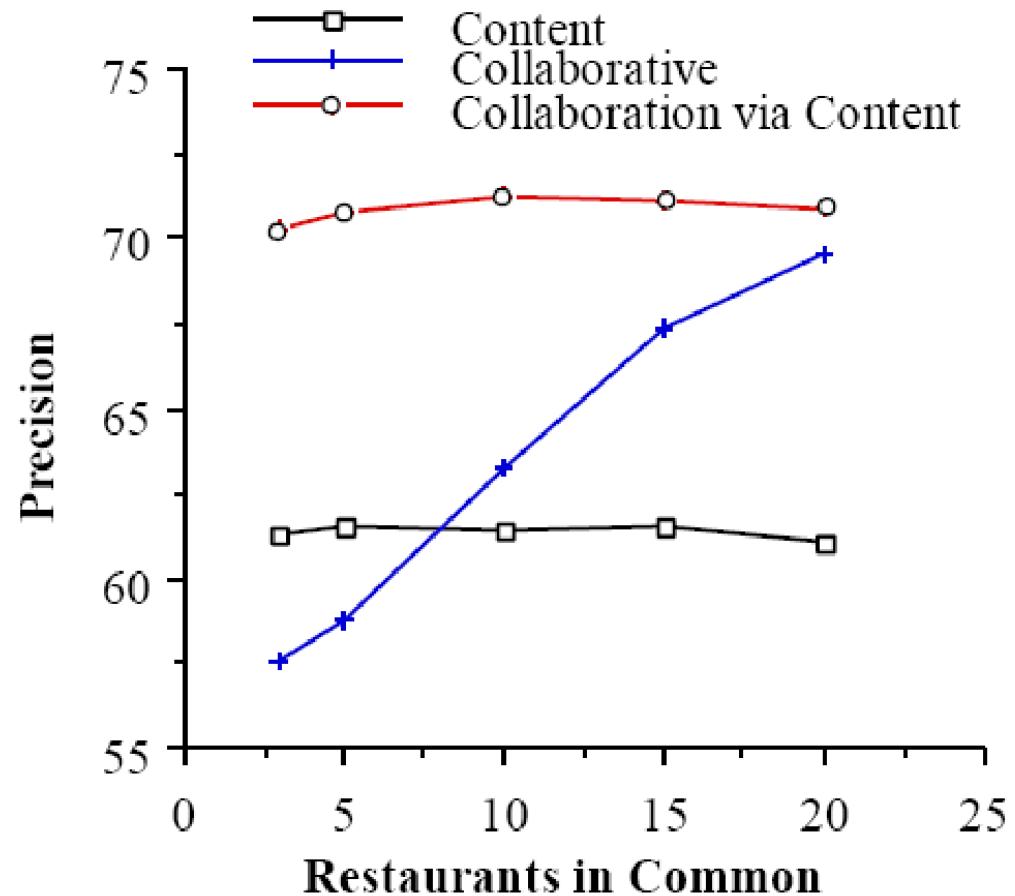
# Index

1. The Recommender Problem
2. “Traditional” Methods
  - 2.1. Collaborative Filtering
  - 2.2. Content-based Recommendations
  - 2.3. Hybrid Approaches
3. Beyond Traditional Methods
  - 3.1. Learning to Rank
  - 3.2. Similarity
  - 3.3. Deep Learning
  - 3.4. Social Recommendations
  - 3.5. Page Optimization
  - 3.6. Tensor Factorization and Factorization Machines
  - 3.7. MAB Explore/Exploit
4. References

## 2.3 Hybrid Approaches

# Comparison of methods (FAB system)

- Content-based recommendation with Bayesian classifier
- Collaborative is standard using Pearson correlation
- Collaboration via content uses the content-based user profiles



Averaged on 44 users

Precision computed in top 3 recommendations

# Hybridization Methods

## Hybridization Method

Weighted

## Description

Outputs from several techniques (in the form of scores or votes) are combined with different degrees of importance to offer final recommendations

Switching

Depending on situation, the system changes from one technique to another

Mixed

Recommendations from several techniques are presented at the same time

Feature combination

Features from different recommendation sources are combined as input to a single technique

Cascade

The output from one technique is used as input of another that refines the result

Feature augmentation

The output from one technique is used as input features to another

Meta-level

The model learned by one recommender is used as input to another



# Index

1. The Recommender Problem
2. “Traditional” Methods
  - 2.1. Collaborative Filtering
  - 2.2. Content-based Recommendations
  - 2.3. Hybrid Approaches
3. Beyond Traditional Methods
  - 3.1. Learning to Rank
  - 3.2. Similarity
  - 3.3. Deep Learning
  - 3.4. Social Recommendations
  - 3.5. Page Optimization
  - 3.6. Tensor Factorization and Factorization Machines
  - 3.7. MAB Explore/Exploit
4. References

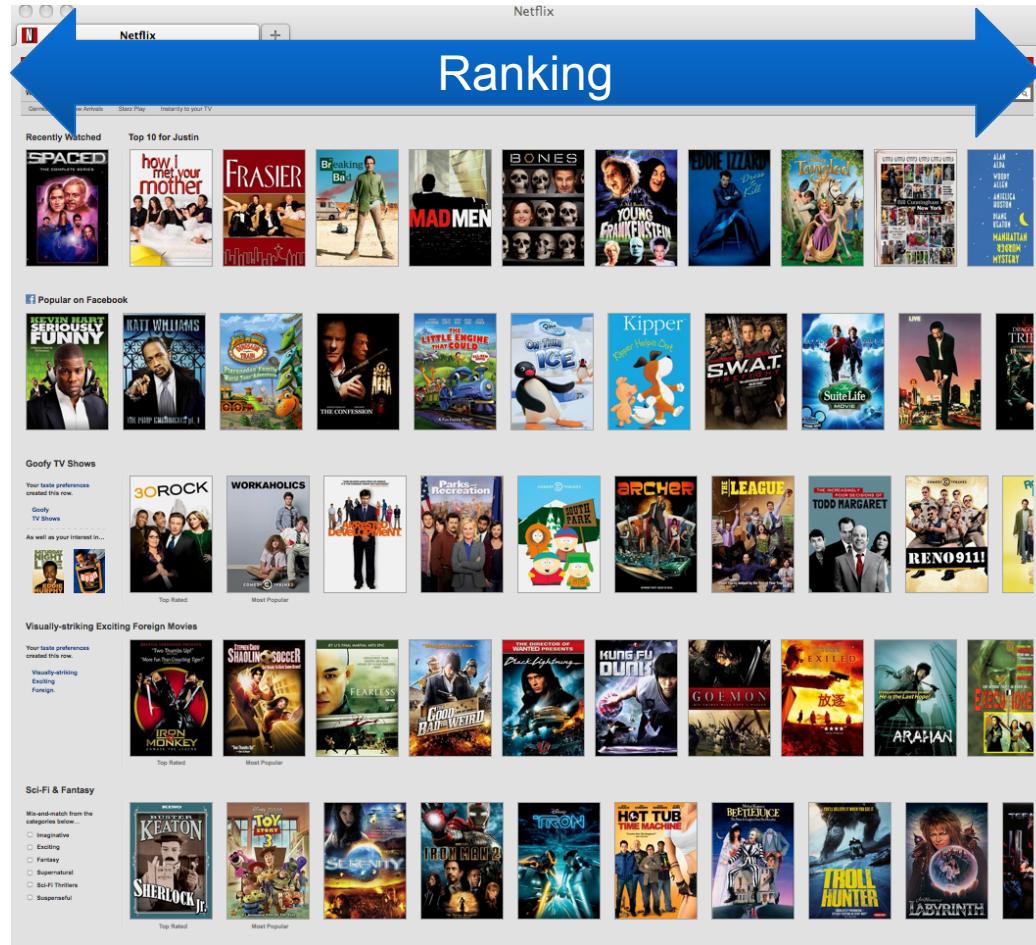
# 3. Beyond traditional approaches to Recommendation



# 3.1 Ranking

# Ranking

Key algorithm, sorts titles in most contexts



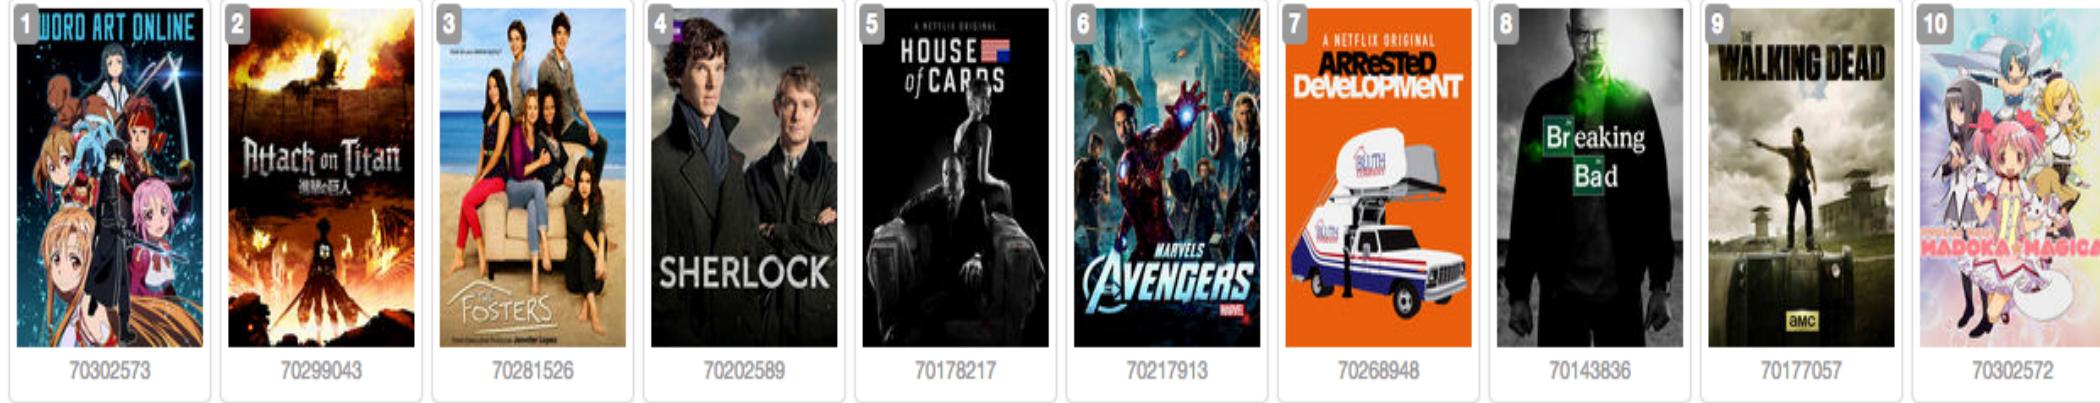
NETFLIX

Xavier Amatriain – October 2014 – Recsys

# Ranking

- Most recommendations are presented in a sorted list
- Recommendation can be understood as a ranking problem
- Popularity is the obvious baseline
- Ratings prediction is a clear secondary data input that allows for personalization
- Many other features can be added

# Ranking by ratings



4.7      4.6      4.5      4.5      4.5      4.5      4.5      4.5      4.5      4.5



Niche titles



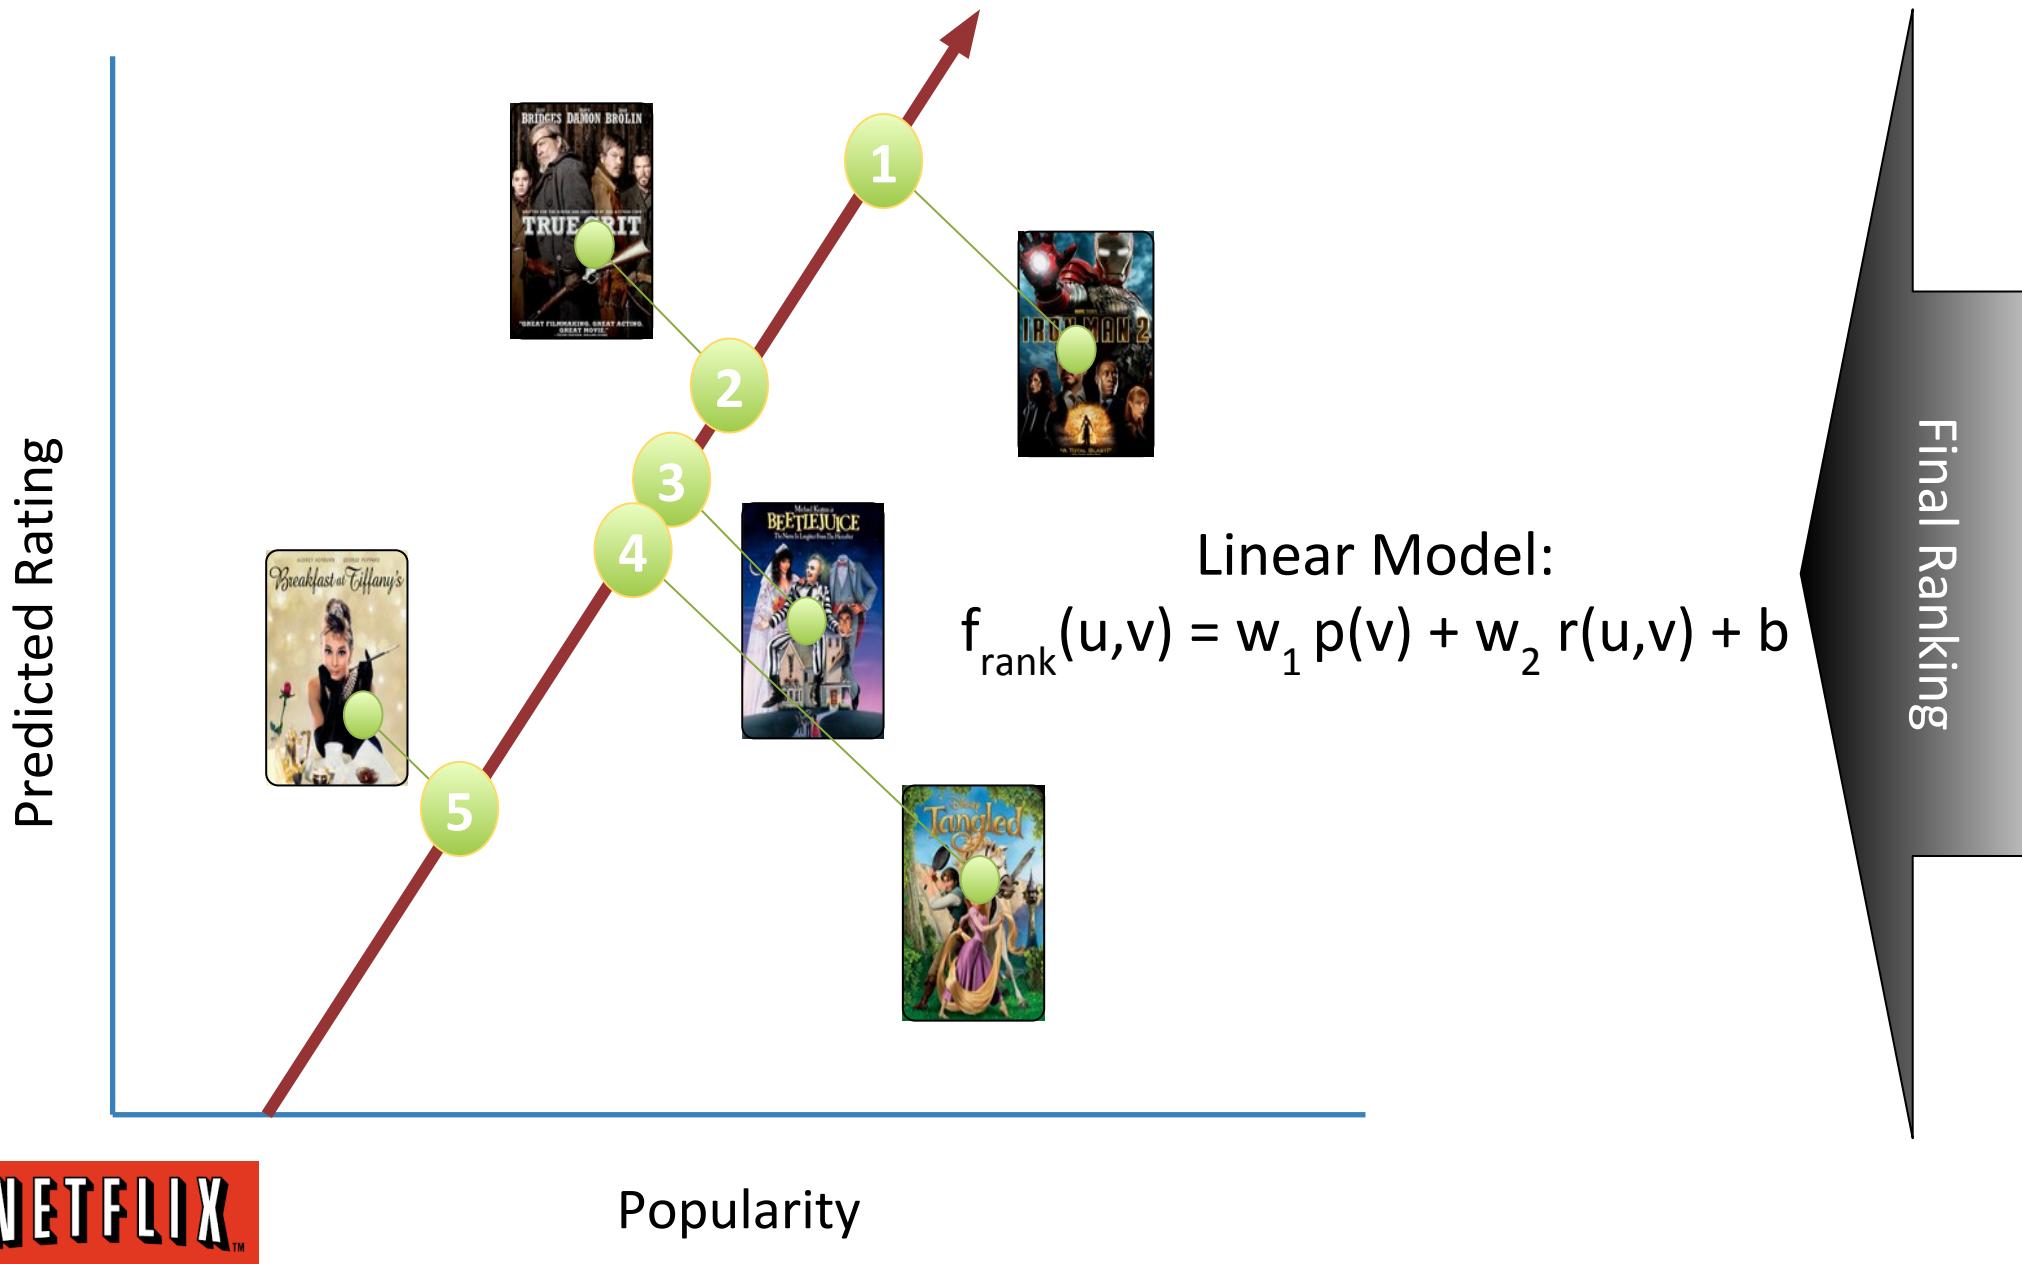
High average ratings... by those who would watch it



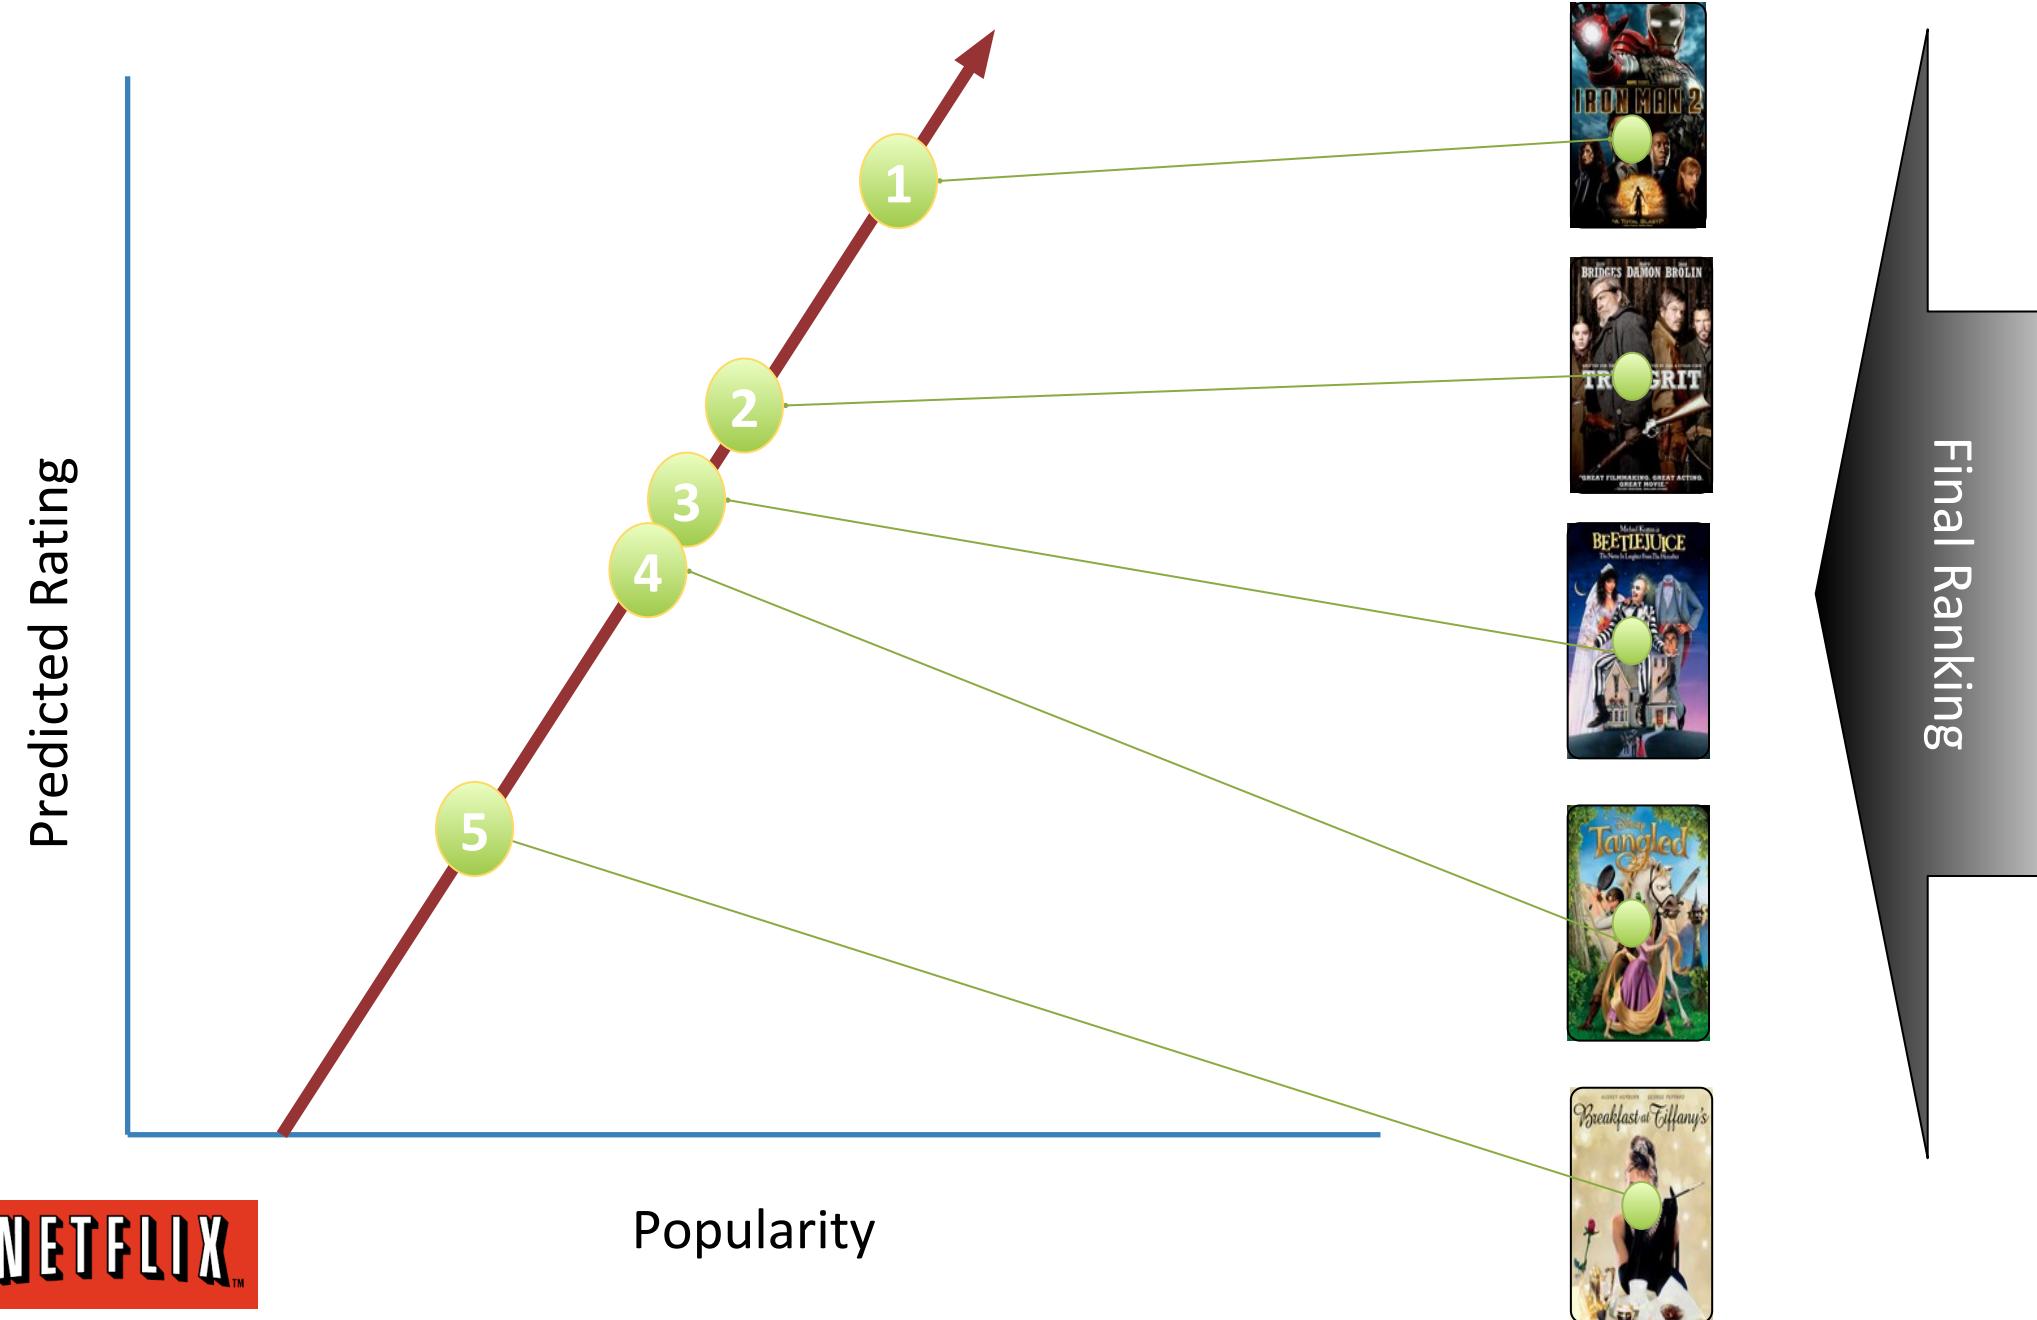
**NETFLIX**

Xavier Amatriain – October 2014 – Recsys

# Example: Two features, linear model



# Example: Two features, linear model



# Learning to rank

- Machine learning problem: goal is to construct ranking model from training data
- Training data can be a partial order or binary judgments (relevant/not relevant).
- Resulting order of the items typically induced from a numerical score
- Learning to rank is a key element for personalization
- You can treat the problem as a standard supervised classification problem

# Learning to rank - Metrics

- Quality of ranking measured using metrics as
  - Normalized Discounted Cumulative Gain
  - Mean Reciprocal Rank (MRR)
  - Fraction of Concordant Pairs (FCP)
  - Others...
- But, it is hard to optimize machine-learned models directly on these measures (e.g. non-differentiable)
- Recent research on models that directly optimize ranking measures

# Learning to rank - Approaches

## 1. Pointwise

- Ranking function minimizes loss function defined on individual relevance judgment
- Ranking score based on regression or classification
- Ordinal regression, Logistic regression, SVM, GBDT, ...

## 2. Pairwise

- Loss function is defined on pair-wise preferences
- Goal: minimize number of inversions in ranking
- Ranking problem is then transformed into the binary classification problem
- RankSVM, RankBoost, RankNet, FRank...

# Learning to rank - Approaches

## 3. Listwise

- Indirect Loss Function
  - RankCosine: similarity between ranking list and ground truth as loss function
  - ListNet: KL-divergence as loss function by defining a probability distribution
  - Problem: optimization of listwise loss function may not optimize IR metrics
- Directly optimizing IR metric (difficult since they are not differentiable)
  - Genetic Programming or Simulated Annealing
  - LambdaMart weights pairwise errors in RankNet by IR metric
  - Gradient descent on smoothed version of objective function (e.g. CLiMF or TFMAP)
  - SVM-MAP relaxes MAP metric by adding to SVM constraints
  - AdaRank uses boosting to optimize NDCG

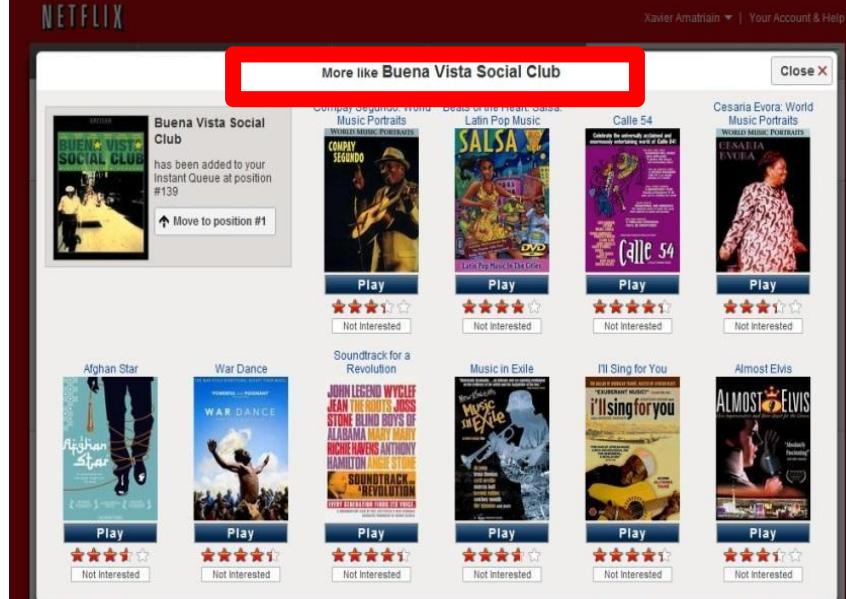
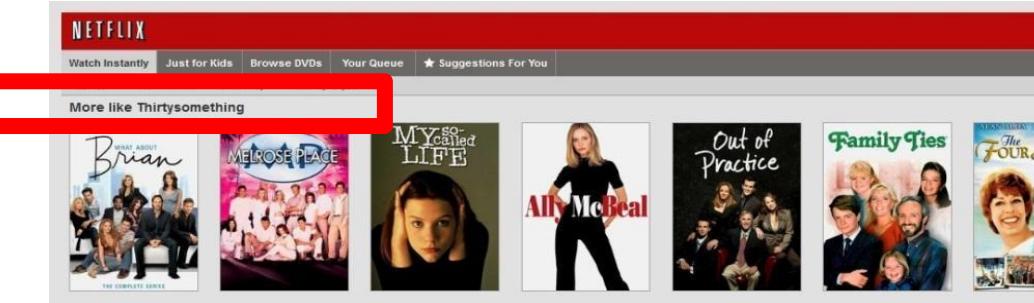
# Index

1. The Recommender Problem
2. “Traditional” Methods
  - 2.1. Collaborative Filtering
  - 2.2. Content-based Recommendations
  - 2.3. Hybrid Approaches
3. Beyond Traditional Methods
  - 3.1. Learning to Rank
  - 3.2. Similarity
  - 3.3. Deep Learning
  - 3.4. Social Recommendations
  - 3.5. Page Optimization
  - 3.6. Tensor Factorization and Factorization Machines
  - 3.7. MAB Explore/Exploit
4. References

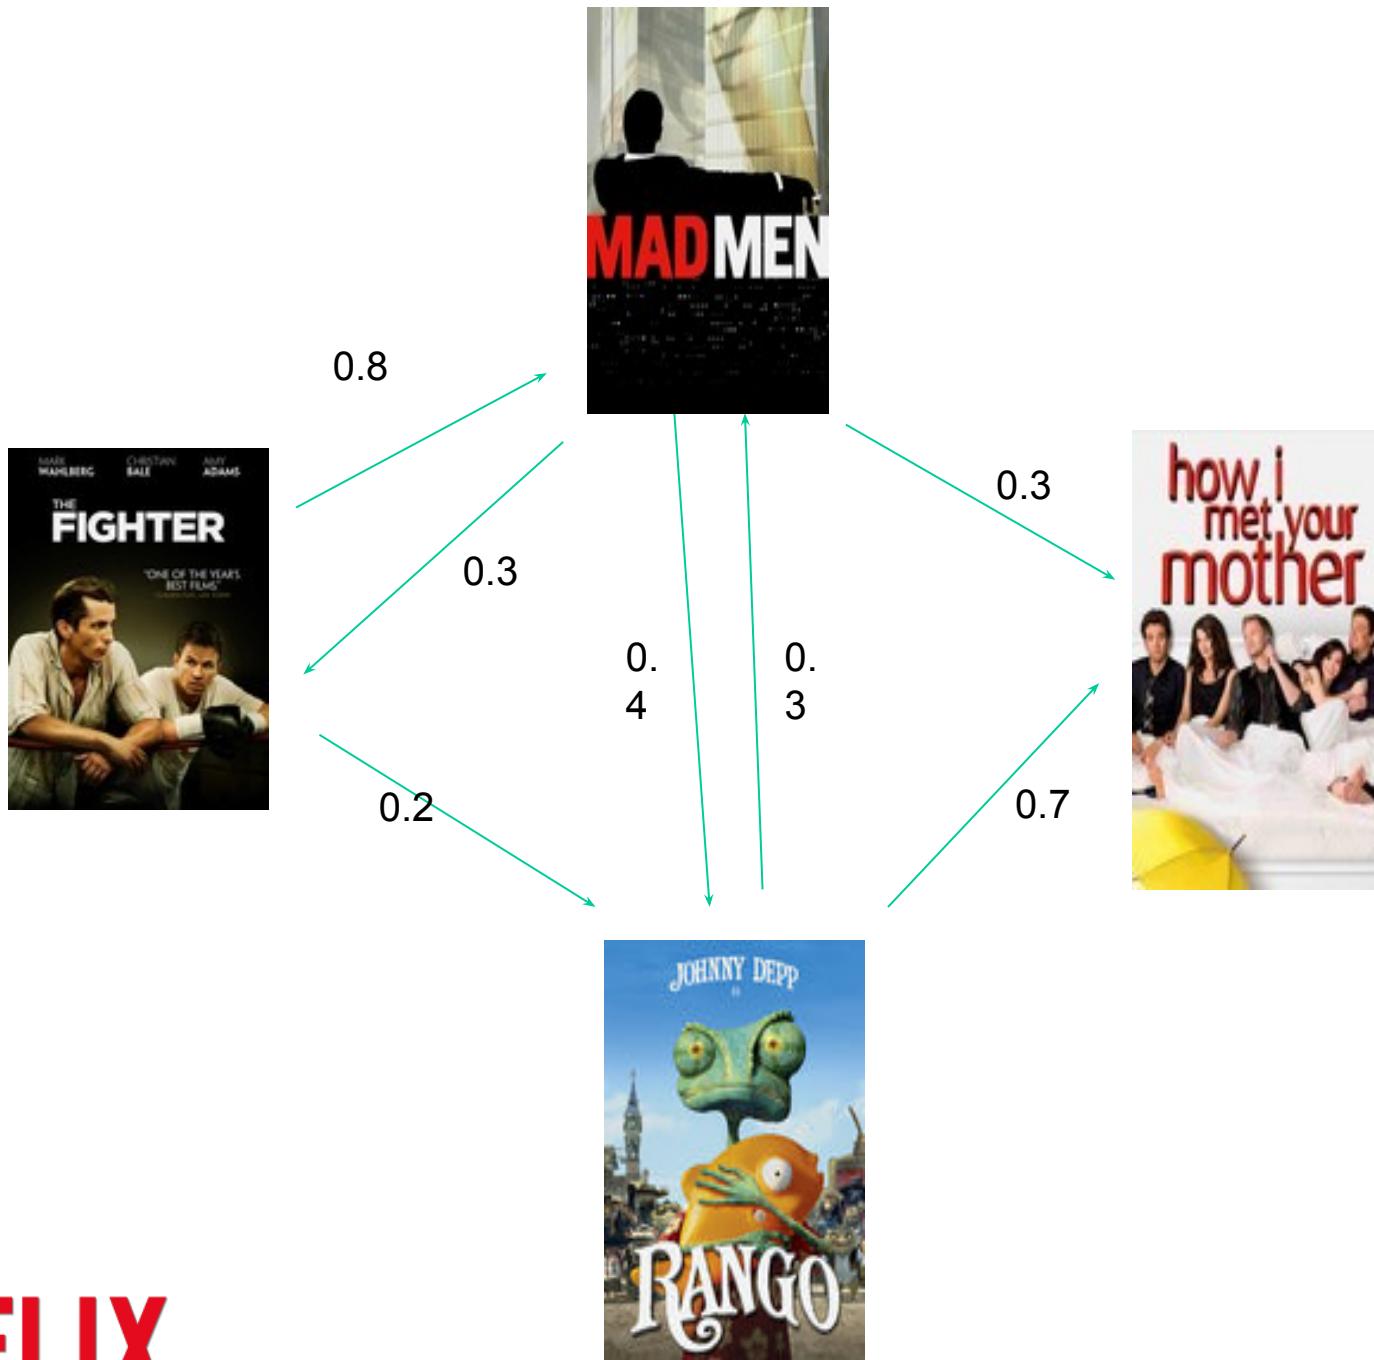
## 3.2 Similarity as Recommendation

# Similar

- Displayed in many different contexts
  - In response to user actions/context (search, queue add...)
  - More like... rows



# Graph-based similarities



NETFLIX

# Example of graph-based similarity: SimRank

- SimRank (Jeh & Widom, 02): “two objects are similar if they are referenced by similar objects.”

$$s(a, b) = \frac{C}{|I(a)||I(b)|} \sum_{i=1}^{|I(a)|} \sum_{j=1}^{|I(b)|} s(I_i(a), I_j(b))$$

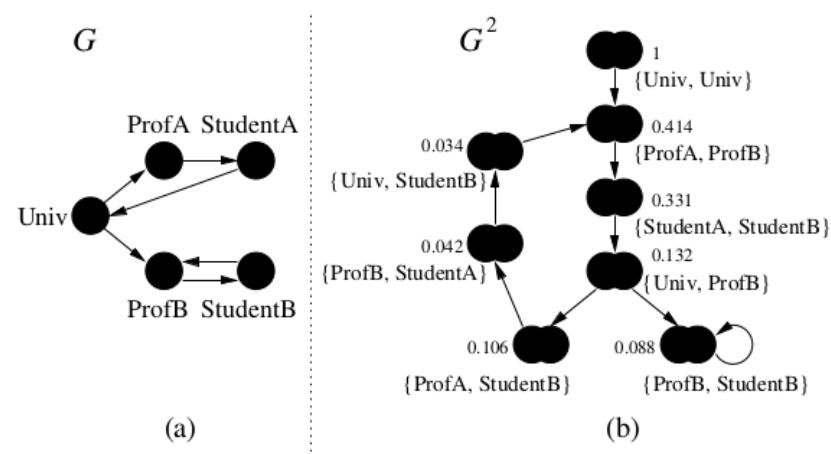


Figure 1: A small Web graph  $G$  and simplified node-pairs graph  $G^2$ . SimRank scores using parameter  $C = 0.8$  are shown for nodes in  $G^2$ .

# Similarity ensembles

- Similarity can refer to different dimensions
  - Similar in metadata/tags
  - Similar in user play behavior
  - Similar in user rating behavior
  - ...
- Combine them using an ensemble
  - Weights are learned using regression over existing response
  - Or... some MAB explore/exploit approach
- The final concept of “similarity” responds to what users vote as similar

# Index

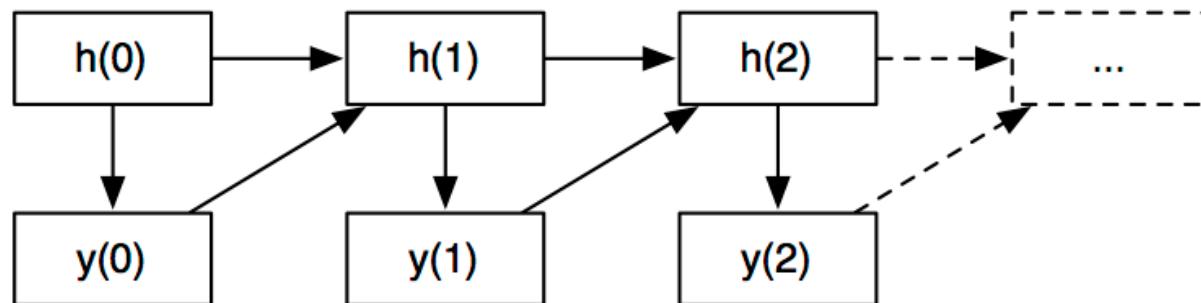
1. The Recommender Problem
2. “Traditional” Methods
  - 2.1. Collaborative Filtering
  - 2.2. Content-based Recommendations
  - 2.3. Hybrid Approaches
3. Beyond Traditional Methods
  - 3.1. Learning to Rank
  - 3.2. Similarity
  - 3.3. Deep Learning
  - 3.4. Social Recommendations
  - 3.5. Page Optimization
  - 3.6. Tensor Factorization and Factorization Machines
  - 3.7. MAB Explore/Exploit
4. References

# 3.3 Deep Learning for Recommendation



# Deep Learning for Collaborative Filtering

- Spotify uses Recurrent Network (<http://erikbern.com/?p=589>)



- In order to predict the next track or movie a user is going to watch, we need to define a distribution  $P(y_i|h_i)$ 
  - If we choose Softmax as it is common practice, we get:
$$P(y_i|h_i) = \frac{\exp(h_t^T a_j)}{\sum_k \exp(h_t^T a_k)}$$
    - Problem: denominator is very expensive to compute
    - Solution: build tree that implements hierarchical softmax
- More details on the blogpost



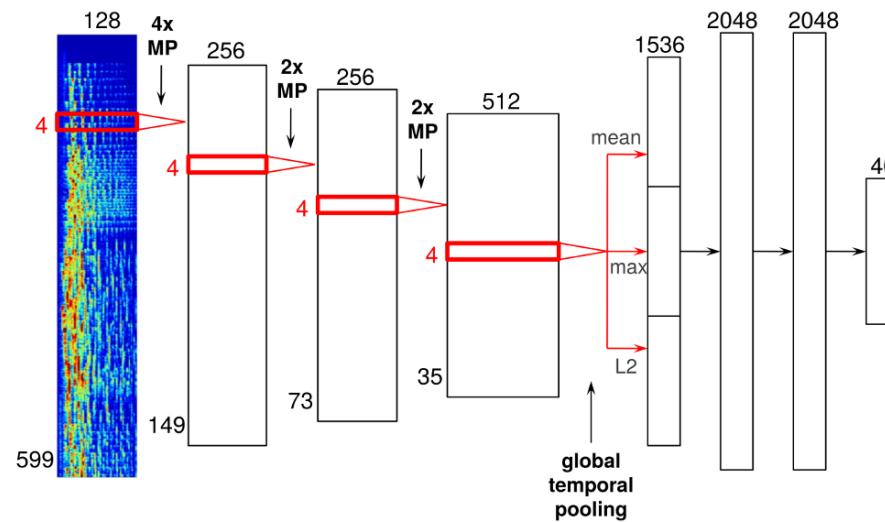
# Deep Learning for Content-based Recommendations

- Another application of Deep Learning to recommendations also from Spotify
  - <http://benanne.github.io/2014/08/05/spotify-cnns.html> also *Deep content-based music recommendation, Aäron van den Oord, Sander Dieleman and Benjamin Schrauwen, NIPS 2013*
- Application to coldstart new titles when very little CF information is available
- Using mel-spectrograms from the audio signal as input
- Training the deep neural network to predict 40 latent factors coming from Spotify's CF solution



# Deep Learning for Content-based Recommendations

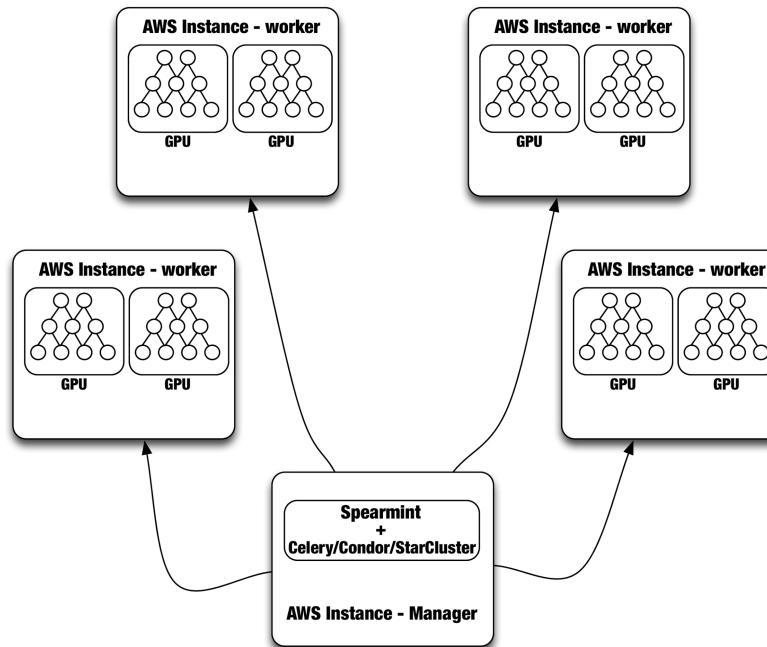
- Network architecture made of 4 convolutional layers + 4 fully connected dense layers



- One dimensional convolutional layers using RELUs (Rectified Linear Units) with activation  $\max(0,x)$
- **Max-pooling operations** between convolutional layers to downsample intermediate representations in time, and add time invariance
- **Global temporal pooling layer** after last convolutional layer: pools across entire time axis, computing statistics of the learned features across time:: mean, maximum and L2-norm
- Globally pooled features are fed into a series of **fully-connected layers** with 2048 RELUs

# ANN Training over GPUS and AWS

- How did we implement our ANN solution at Netflix?
  - Level 1 distribution: machines over different AWS regions
  - Level 2 distribution: machines in AWS and same AWS region
    - Use coordination tools
      - Spearmint or similar for parameter optimization
      - Condor, StarCluster, Mesos... for distributed cluster coordination
  - Level 3 parallelization: highly optimized parallel CUDA code on GPUs



<http://techblog.netflix.com/2014/02/distributed-neural-networks-with-gpus.html>



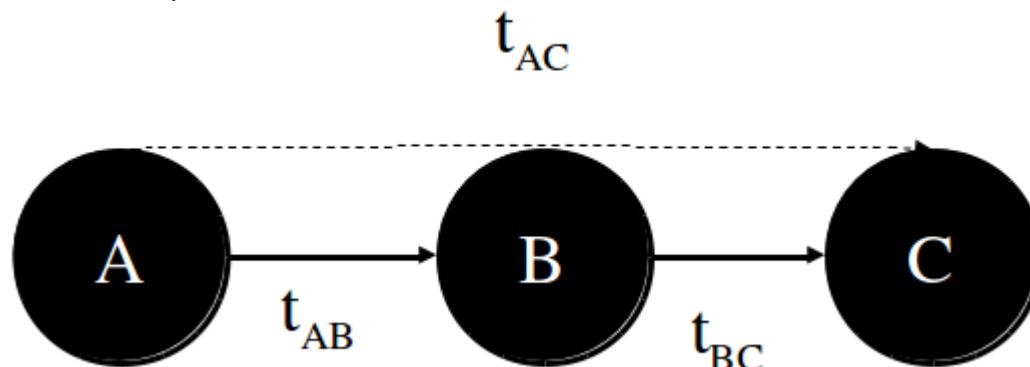
# Index

1. The Recommender Problem
2. “Traditional” Methods
  - 2.1. Collaborative Filtering
  - 2.2. Content-based Recommendations
  - 2.3. Hybrid Approaches
3. Beyond Traditional Methods
  - 3.1. Learning to Rank
  - 3.2. Similarity
  - 3.3. Deep Learning
  - 3.4. Social Recommendations
  - 3.5. Page Optimization
  - 3.6. Tensor Factorization and Factorization Machines
  - 3.7. MAB Explore/Exploit
4. References

## 3.4 Social Recommendations

# Social and Trust-based recommenders

- A social recommender system recommends items that are “popular” in the social proximity of the user.
- Social proximity = trust (can also be topic-specific)
- Given two individuals - the source (node A) and *sink* (node C) - derive how much the source should trust the sink.
- Algorithms
  - Advogato (Levien)
  - Appleseed (Ziegler and Lausen)
  - MoleTrust (Massa and Avesani)
  - TidalTrust (Golbeck)



# Other ways to use Social

- Social connections can be used in combination with other approaches
- In particular, “friendships” can be fed into collaborative filtering methods in different ways
  - replace or modify user-user “similarity” by using social network information
  - use social connection as a part of the ML objective function as regularizer



# Demographic Methods

- Aim to categorize the user based on personal attributes and make recommendation based on demographic classes
- Demographic groups can come from marketing research – hence experts decided how to model the users
- Demographic techniques form people-to-people correlations

# Index

1. The Recommender Problem
2. “Traditional” Methods
  - 2.1. Collaborative Filtering
  - 2.2. Content-based Recommendations
  - 2.3. Hybrid Approaches
3. Beyond Traditional Methods
  - 3.1. Learning to Rank
  - 3.2. Similarity
  - 3.3. Deep Learning
  - 3.4. Social Recommendations
  - 3.5. Page Optimization
  - 3.6. Tensor Factorization and Factorization Machines
  - 3.7. MAB Explore/Exploit
4. References

## 3.5. Page Optimization

# Page Composition

NETFLIX

Watch Instantly · Just for Kids · Taste Profile · DVDs · DVD Queue

Recently Watched · My List · See All

Popular on Netflix

Critically-acclaimed Movies

New Releases

Kids' TV

Because you watched Revolution

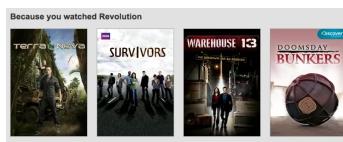
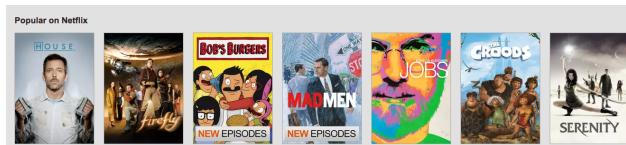
Because you watched Stargate

Top Picks for Justin

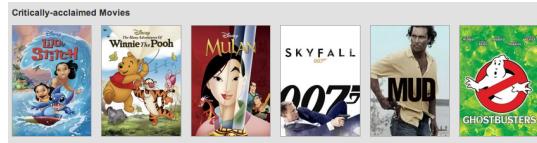
Children & Family Movies

# Page Composition

10,000s  
of  
possible  
rows



...



Variable number of  
possible videos per  
row (up to  
thousands)

1 personalized  
page



10-40  
rows

per  
device



# Page Composition

The image shows two side-by-side web pages. On the left is a Google search results page for the query "Louvre 2006 donation". The results include several links related to the Louvre Museum's 2006 restoration efforts, including news articles and reviews. On the right is a Wikipedia article titled "Nexus 5". The page contains detailed information about the smartphone, including its hardware specifications, software features, and release history. A purple arrow points from the bottom of the Google search results towards the top of the Nexus 5 page.

Google search results for "Louvre 2006 donation":

- Donating online | Louvre Museum | Paris - Musée du Louvre
- Support the Louvre | Louvre Museum | Paris - Musée du Louvre
- Louvre asks for donations to restore the Nile of Samothrace - Euronews
- Louvre Museum
- Louvre Gets \$20 Million for New Islamic Wing - New York Times
- On a Mission to Looosen Up the Louvre - The New York Times
- A Francisco Millet painting donated to the Louvre - The Art Tribune
- Several French drawings donated to the Louvre on condition - The Art Tribune
- The Louvre - Wikipedia, the free encyclopedia
- Did Michelangelo Make Crucifix Donated to Louvre? - YouTube

Wikipedia article on the Nexus 5:

**Nexus 5**

The Nexus 5 is a smartphone co-developed by Google and LG Electronics that runs the Android operating system. The successor to the Nexus 4, the device is the fifth smartphone in the Google Nexus series, a family of Android consumer devices marketed by Google and built by an original equipment manufacturer partner. The Nexus 5 was unveiled on 31 October 2013, and released the same day for online purchase on Google Play, in selected countries.

The Nexus 5's hardware is similar to that of the LG G2, with a Snapdragon 800 system-on-chip (SoC), and a 4.96-inch 1080p display. The Nexus 5 is also the first device to feature version 4.4 of Android.

**Release [edit]**

The Nexus 5 was initially released for ordering at Google Play Store on October 31, 2013, in 16 GB and 32 GB versions.

**Specifications [edit]**

**Hardware [edit]**

Its exterior is made from a polycarbonate shell with similarities to the new Nexus 7, unlike its predecessor, which used a glass-based construction.

It's hardware contains similarities to the LG G2; it is powered by a 2.26 GHz quad-core Snapdragon 800 processor with 2 GB of RAM, either 16 or 32 GB of internal storage, and a 2300 mAh battery. The Nexus 5 uses a 4.96-inch (measured as 5-inch) 1080p IPS 1080p display, and includes an 8-megapixel rear-facing camera with optical image stabilization (OIS). The Nexus 5 supports LTE networks where available, unlike the Nexus 4 which unofficially supported LTE on AT&T Band 4 only with a hidden software option, but was not formally approved or marketed for any LTE use.

There are two variants of the Nexus 5; one is specific to North America (LG-D820), and the other is designed for the rest of the world (LG-D821). The differences between these two variants are in supported cellular frequency bands, see the infobox on the right for more details.

Like its predecessor, the Nexus 5 does not have a microSD card slot, while it features a multi-color LED notification light. Despite the fact there is a pair of speakers present on the lower edge of the Nexus 5, there is only one speaker, one grille is for a speaker, and another is for a microphone.

Notable new hardware features also include two new composite sensors: a step detector and a step counter. These new sensors allow applications to easily track steps when the user is walking, running, or climbing stairs. Both sensors are implemented in hardware for low power consumption.

**Software [edit]**

Nexus 5 is the first Android device to ship with Android 4.4 "KitKat", which has a refreshed interface, improved performance, improved multitasking (such as the ability to emulate a smart card), a new "HDR+" camera shooting mode, native printing functionality, a screen rotation lock, and other new and improved functionality.

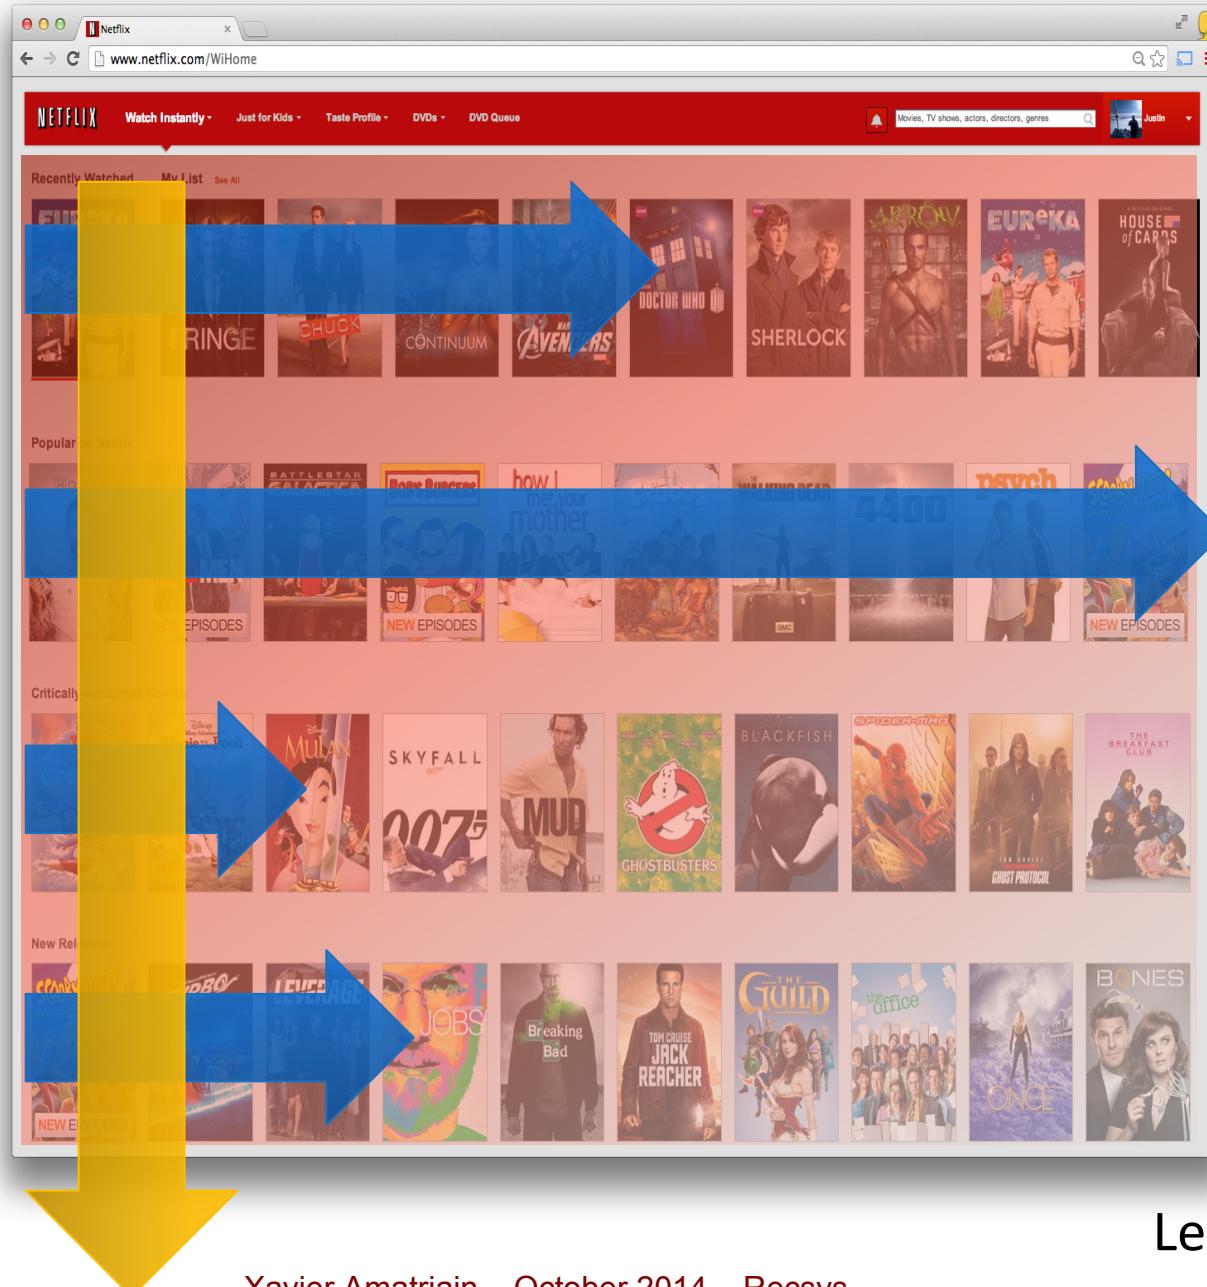
Nexus 5 ships with Google Experience Launcher (GEL), a redesigned home screen which allows users to access Google Now on a dedicated page, and allows voice search to be activated on the home screen with a voice command. Unlike other features of Android 4.4, GEL is not backwards compatible; it is backwards incompatible with the Google Search application as of November 2013. GEL

From “Modeling User Attention and Interaction on the Web” 2014 - PhD Thesis by Dmitry Lagun (Emory U.)



# User Attention Modeling

More likely to see



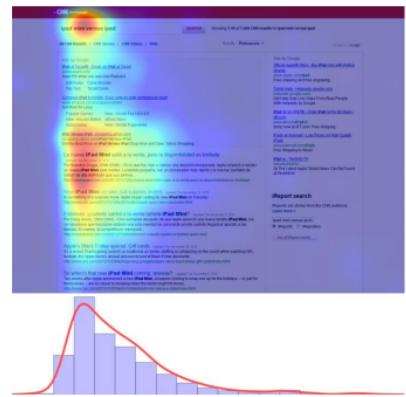
NETFLIX

# User Attention Modeling

Web Search (Google) Social Network (Twitter)



News (CNN)



Shopping (Amazon)



From “Modeling User Attention and Interaction on the Web” 2014 - PhD Thesis by Dmitry Lagun (Emory U.)



Xavier Amatriain – October 2014 – Recsys

# Page Composition

Accurate vs. Diverse  
Discovery vs. Continuation  
Depth vs. Coverage  
Freshness vs. Stability  
Recommendations vs. Tasks

- To put things together we need to combine different elements
  - Navigational/Attention Model
  - Personalized Relevance Model
  - Diversity Model

## Fair and Balanced: Learning to Present News Stories

Amr Ahmed<sup>\*1</sup>, Choon Hui Teo<sup>\*1</sup>, S.V.N. Vishwanathan<sup>2</sup>, Alex Smola<sup>1</sup>

\* Co-first authors.

<sup>1</sup>Yahoo! Research, Santa Clara, CA 95053, USA

<sup>2</sup>Purdue University, West Lafayette, IN 47907, USA

{amahmed,choonhui,smola}@yahoo-inc.com, vishy@stat.purdue.edu



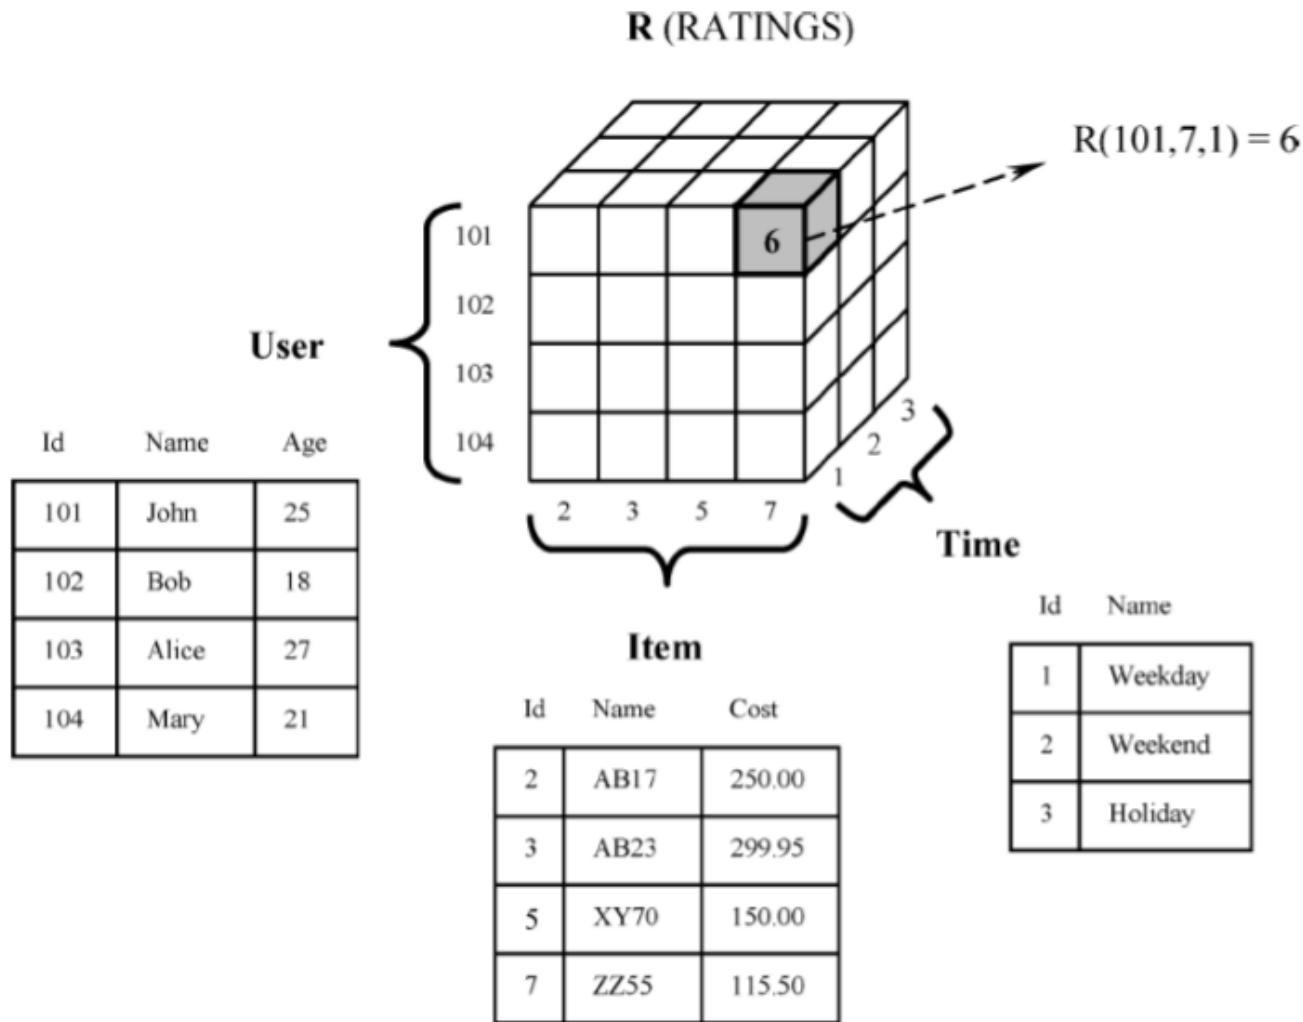
# Index

1. The Recommender Problem
2. “Traditional” Methods
  - 2.1. Collaborative Filtering
  - 2.2. Content-based Recommendations
  - 2.3. Hybrid Approaches
3. Beyond Traditional Methods
  - 3.1. Learning to Rank
  - 3.2. Similarity
  - 3.3. Deep Learning
  - 3.4. Social Recommendations
  - 3.5. Page Optimization
  - 3.6. Tensor Factorization and Factorization Machines
  - 3.7. MAB Explore/Exploit
4. References

# 3.6 Tensor Factorization & Factorization Machines



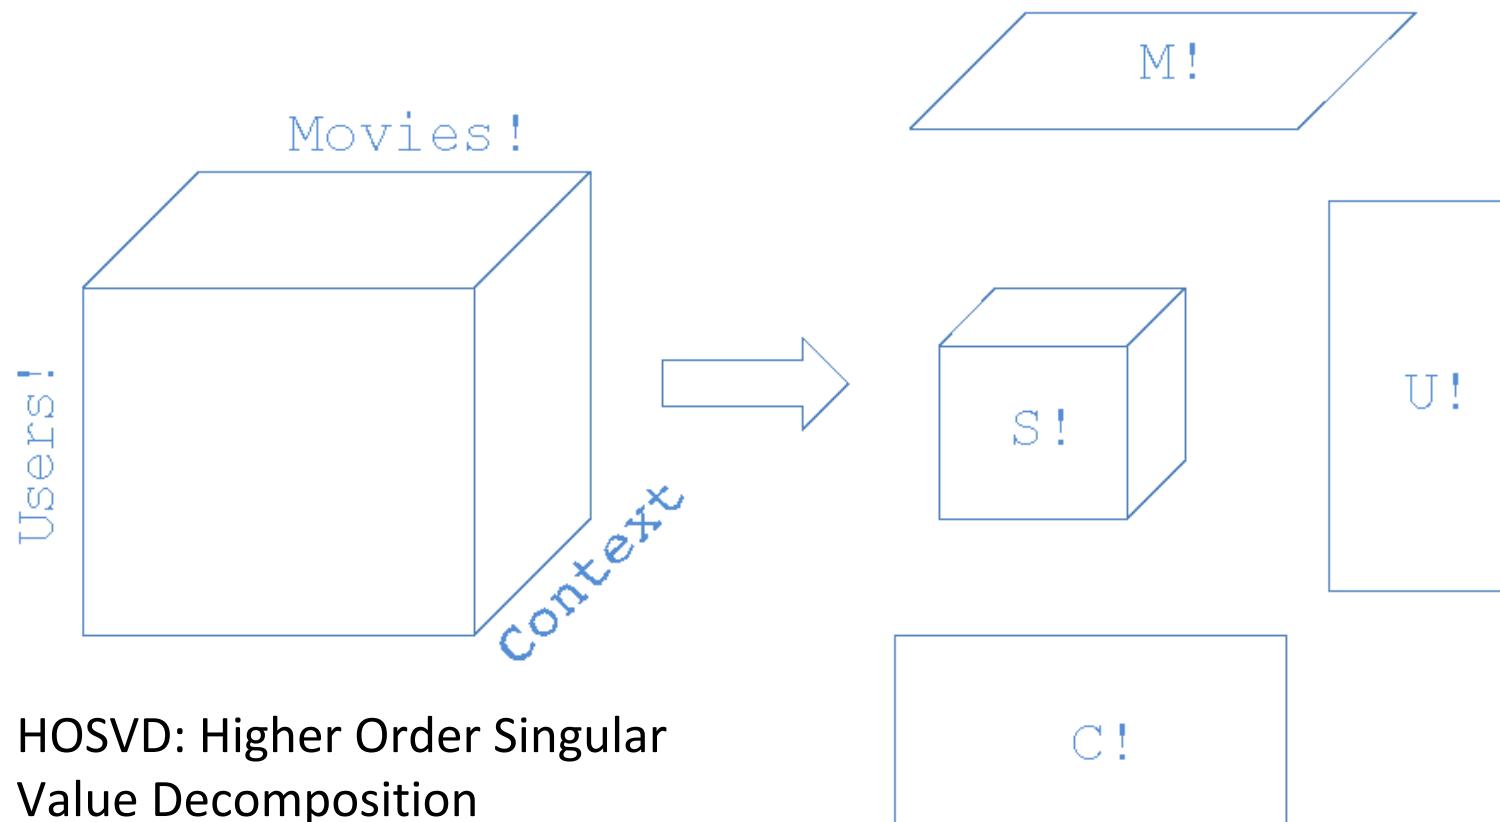
# N-dimensional model



[Adomavicius et al., 2005]



# Tensor Factorization



HOSVD: Higher Order Singular  
Value Decomposition

$$U \in \mathbb{R}^{n \times d_U}, M \in \mathbb{R}^{m \times d_M} \text{ and } C \in \mathbb{R}^{c \times d_C}$$

$$S \in \mathbb{R}^{d_U \times d_M \times d_C}$$

**HOSVD  
Model**

$$R[U, M, C, S] := L(F, Y) + \Omega[U, M, C] + \Omega[S]$$



# Factorization Machines

- Generalization of regularized matrix (and tensor) factorization approaches combined with linear (or logistic) regression
- Problem: Each new adaptation of matrix or tensor factorization requires deriving new learning algorithms
- Combines “generality of machine learning/regression with quality of factorization models”



## Gradient boosting factorization machines

Authors: [Chen Cheng](#) The Chinese University of Hong Kong, Hong Kong, Hong Kong  
[Fen Xia](#) Baidu Inc., Beijing, China  
[Tong Zhang](#) Baidu Inc., Beijing, China  
[Irwin King](#) The Chinese University of Hong Kong, Hong Kong, Hong Kong  
[Michael R. Lyu](#) The Chinese University of Hong Kong, Hong Kong, Hong Kong



20



- Downloads (6)
- Downloads (12)
- Downloads (cum)
- Citation Count

### Published in:



• Proceeding  
[RecSys '14](#) Proceedings of the 8th ACM Conference on Recommender systems  
Pages 265-272  
ACM New York, NY, USA ©2014  
[table of contents](#) ISBN: 978-1-4503-2668-1 doi:>[10.1145/2645710.2645730](https://doi.org/10.1145/2645710.2645730)



## 'Free lunch' enhancement for collaborative filtering with factorization machines

Authors: [Babak Loni](#) Delft University of Technology, Delft, Netherlands  
[Alan Said](#) Delft University of Technology, Delft, Netherlands  
[Martha Larson](#) Delft University of Technology, Delft, Netherlands  
[Alan Hanjalic](#) Delft University of Technology, Delft, Netherlands



Bibliometrics

- Downloads (6 We)
- Downloads (12 M)
- Downloads (cum)
- Citation Count: 0

### Published in:



• Proceeding  
[RecSys '14](#) Proceedings of the 8th ACM Conference on Recommender systems  
Pages 281-284  
ACM New York, NY, USA ©2014  
[table of contents](#) ISBN: 978-1-4503-2668-1 doi:>[10.1145/2645710.2645771](https://doi.org/10.1145/2645710.2645771)



# Factorization Machines

- Each feature gets a weight value and a factor vector
  - $O(dk)$  parameters

$$b \in \mathbb{R}, \mathbf{w} \in \mathbb{R}^d, \mathbf{V} \in \mathbb{R}^{d \times k}$$

- Model equation:

$$\begin{aligned} f(\mathbf{x}) &= b + \sum_{i=1}^d w_i x_i + \sum_{i=1}^d \sum_{j=i+1}^d x_i x_j \mathbf{v}_i^T \mathbf{v}_j & O(d^2) \\ &= b + \sum_{i=1}^d w_i x_i + \frac{1}{2} \sum_{f=1}^k \left( \left( \sum_{i=1}^d x_i v_{i,f} \right)^2 - \sum_{i=1}^d x_i^2 v_{i,f}^2 \right) & O(kd) \end{aligned}$$



# Factorization Machines

- Two categorical variables ( $u, i$ ) encoded as real values:

Feature vector $\mathbf{x}$									
$\mathbf{x}^{(1)}$	1	0	0	...	1	0	0	0	...
$\mathbf{x}^{(2)}$	1	0	0	...	0	1	0	0	...
$\mathbf{x}^{(3)}$	1	0	0	...	0	0	1	0	...
$\mathbf{x}^{(4)}$	0	1	0	...	0	0	1	0	...
$\mathbf{x}^{(5)}$	0	1	0	...	0	0	0	1	...
$\mathbf{x}^{(6)}$	0	0	1	...	1	0	0	0	...
$\mathbf{x}^{(7)}$	0	0	1	...	0	0	1	0	...
	A	B	C	...	TI	NH	SW	ST	...
	User				Movie				

- FM becomes identical to MF with biases:

$$f(\mathbf{x}) = b + w_u + w_i + \mathbf{v}_u^T \mathbf{v}_i$$

From Rendle (2012) KDD Tutorial



# Factorization Machines

- Makes it easy to add a time signal

Feature vector $\mathbf{x}$										
$\mathbf{x}^{(1)}$	1	0	0	...	1	0	0	0	...	0.2
$\mathbf{x}^{(2)}$	1	0	0	...	0	1	0	0	...	0.6
$\mathbf{x}^{(3)}$	1	0	0	...	0	0	1	0	...	0.61
$\mathbf{x}^{(4)}$	0	1	0	...	0	0	1	0	...	0.3
$\mathbf{x}^{(5)}$	0	1	0	...	0	0	0	1	...	0.5
$\mathbf{x}^{(6)}$	0	0	1	...	1	0	0	0	...	0.1
$\mathbf{x}^{(7)}$	0	0	1	...	0	0	1	0	...	0.8
	A	B	C	...	TI	NH	SW	ST	...	Time
	User				Movie				Time	

- Equivalent equation:

$$f(\mathbf{x}) = b + w_u + w_i + x_t w_t + \mathbf{v}_u^T \mathbf{v}_i + x_t \mathbf{v}_u^T \mathbf{v}_t + x_t \mathbf{v}_i^T \mathbf{v}_t$$

From Rendle (2012) KDD Tutorial



# Factorization Machines

- L2 regularized
  - Regression: Optimize RMSE
  - Classification: Optimize logistic log-likelihood
  - Ranking: Optimize scores
- Can be trained using:
  - SGD
  - Adaptive SGD
  - ALS
  - MCMC

$$\frac{\partial}{\partial \theta} f(\mathbf{x}) = \begin{cases} 1 & \text{if } \theta \text{ is } b \\ x_i & \text{if } \theta \text{ is } w_i \\ x_i \sum_{j=1}^d v_{j,f} x_j - v_{i,f} x_i^2 & \text{if } \theta \text{ is } v_{i,f} \end{cases}$$

Least squares SGD:

$$\theta' = \theta - \eta \left( (f(\mathbf{x}) - y) \frac{\partial}{\partial \theta} f(\mathbf{x}) + \lambda_\theta \theta \right)$$



# Index

1. The Recommender Problem
2. “Traditional” Methods
  - 2.1. Collaborative Filtering
  - 2.2. Content-based Recommendations
  - 2.3. Hybrid Approaches
3. Beyond Traditional Methods
  - 3.1. Learning to Rank
  - 3.2. Similarity
  - 3.3. Deep Learning
  - 3.4. Social Recommendations
  - 3.5. Page Optimization
  - 3.6. Tensor Factorization and Factorization Machines
  - 3.7. MAB Explore/Exploit
4. References

# 3.7 MAB Explore/Exploit



# Explore/Exploit

- One of the key issues when building any kind of personalization algorithm is how to trade off:
  - **Exploitation**: Cashing in on what we know about the user right now
  - **Exploration**: Using the interaction as an opportunity to learn more about the user
- We need to have informed and optimal strategies to drive that tradeoff
  - **Solution**: pick a reasonable set of candidates and show users only “enough” to gather information on them



# Multi-armed Bandits

- Given possible strategies/candidates (slot machines) pick the arm that has the maximum potential of *being good* (minimize regret)



- Naive strategy:  $\epsilon$  – *greedy*
  - Explore with a small probability  $\epsilon$  (e.g. 5%) -> choose an arm at random
  - Exploit with a high probability ( $1 - \epsilon$ ) (e.g. 95%) -> choose the best-known arm so far
- Translation to recommender systems
  - Choose an arm = choose an item/choose an algorithm (MAB testing)

# Multi-armed Bandits

- Better strategies not only take into account the mean of the posterior, but also the variance
- Upper Confidence Bound (UCB)
  - Show item with maximum score
  - Score = Posterior mean +  $\Delta$
  - Where  $\Delta$  can be computed differently depending on the variant of UCB (e.g.  $\propto$  to the number of trials or the measured variance)
- Thompson Sampling
  - Given a posterior distribution  $P(\theta|\mathcal{D}) \propto P(\mathcal{D}|\theta)P(\theta)$
  - Sample on each iteration and choose the action that maximizes the expected reward

# Multi-armed Bandits

## A Contextual-Bandit Approach to Personalized News Article Recommendation

### Explore-Exploit in Top-N Recommender Systems via Gaussian Processes

Hastagiri P Vanchinathan  
ETH Zürich  
hastagiri@inf.ethz.ch

Isidor Nikolic \*  
Microsoft, Zürich  
inikolic@microsoft.com

Andreas Krause  
ETH Zürich  
krausea@ethz.ch

Fabio De Bona  
Google, Zürich  
fdb@google.com

Lihong Li†, Wei Chu†,  
†Yahoo! Labs  
lihong.chuwei@yahoo-inc.com

John Langford‡  
‡Yahoo! Labs  
jl@yahoo-inc.com

Robert E. Schapire+\*  
+Dept of Computer Science  
Princeton University  
schapire@cs.princeton.edu

### Context Adaptation in Interactive Recommender Systems

Negar Hariri  
DePaul University  
Chicago, IL 60604, USA  
nhariri@cs.depaul.edu

Bamshad Mobasher  
DePaul University  
Chicago, IL 60604, USA  
mobasher@cs.depaul.edu

Robin Burke  
DePaul University  
Chicago, IL 60604, USA  
rburke@cs.depaul.edu



LinkedIn

Recommending Items to Users: An Explore Exploit Perspective

Deepak Agarwal, Director Machine Learning and Relevance Science, LinkedIn, USA

CIKM, 2013



# Index

1. The Recommender Problem
2. “Traditional” Methods
  - 2.1. Collaborative Filtering
  - 2.2. Content-based Recommendations
  - 2.3. Hybrid Approaches
3. Beyond Traditional Methods
  - 3.1. Learning to Rank
  - 3.2. Similarity
  - 3.3. Deep Learning
  - 3.4. Social Recommendations
  - 3.5. Page Optimization
  - 3.6. Tensor Factorization and Factorization Machines
  - 3.7. MAB Explore/Exploit
4. References

# 4. References

# References

- "Recommender Systems Handbook." Ricci, Francesco, Lior Rokach, Bracha Shapira, and Paul B. Kantor. (2010).
- "Recommender systems: an introduction". Jannach, Dietmar, et al. Cambridge University Press, 2010.
- "Toward the Next Generation of Recommender Systems: A Survey of the State-of-the-Art and Possible Extensions". G. Adomavicius and A. Tuzhilin. 2005. IEEE Transactions on Knowledge and Data Engineering, 17 (6)
- "Item-based Collaborative Filtering Recommendation Algorithms", B. Sarwar et al. 2001. Proceedings of World Wide Web Conference.
- "Lessons from the Netflix Prize Challenge.". R. M. Bell and Y. Koren. SIGKDD Explor. Newsl., 9(2):75–79, December 2007.
- "Beyond algorithms: An HCI perspective on recommender systems". K. Swearingen and R. Sinha. In ACM SIGIR 2001 Workshop on Recommender Systems
- "Recommender Systems in E-Commerce". J. Ben Schafer et al. ACM Conference on Electronic Commerce. 1999-
- "Introduction to Data Mining", P. Tan et al. Addison Wesley. 2005



# References

- “*Evaluating collaborative filtering recommender systems*”. J. L. Herlocker, J. A. Konstan, L. G. Terveen, and J. T. Riedl. ACM Trans. Inf. Syst., 22(1): 5–53, 2004.
- “*Trust in recommender systems*”. J. O’Donovan and B. Smyth. In Proc. of IUI ’05, 2005.
- “*Content-based recommendation systems*”. M. Pazzani and D. Billsus. In The Adaptive Web, volume 4321. 2007.
- “*Fast context-aware recommendations with factorization machines*”. S. Rendle, Z. Gantner, C. Freudenthaler, and L. Schmidt-Thieme. In Proc. of the 34th ACM SIGIR, 2011.
- “*Restricted Boltzmann machines for collaborative filtering*”. R. Salakhutdinov, A. Mnih, and G. E. Hinton. In Proc of ICML ’07, 2007
- “Learning to rank: From pairwise approach to listwise approach”. Z. Cao and T. Liu. In In Proceedings of the 24th ICML, 2007.
- “*Introduction to Data Mining*”, P. Tan et al. Addison Wesley. 2005

# References

- D. H. Stern, R. Herbrich, and T. Graepel. “*Matchbox: large scale online bayesian recommendations*”. In Proc.of the 18th WWW, 2009.
- Koren Y and J. Sill. “*OrdRec: an ordinal model for predicting personalized item rating distributions*”. In Rec-Sys '11.
- Y. Koren. “*Factorization meets the neighborhood: a multifaceted collaborative filtering model*”. In Proceedings of the 14th ACM SIGKDD, 2008.
- Yifan Hu, Y. Koren, and C. Volinsky. “*Collaborative Filtering for Implicit Feedback Datasets*”. In Proc. Of the 2008 Eighth ICDM
- Y. Shi, A. Karatzoglou, L. Baltrunas, M. Larson, N. Oliver, and A. Hanjalic. “*CLiMF: learning to maximize reciprocal rank with collaborative less-is-more filtering*”. In Proc. of the sixth Recsys, 2012.
- Y. Shi, A. Karatzoglou, L. Baltrunas, M. Larson, A. Hanjalic, and N. Oliver. “*TFMAP: optimizing MAP for top-n context-aware recommendation*”. In Proc. Of the 35th SIGIR, 2012.
- C. Burges. 2010. *From RankNet to LambdaRank to LambdaMART: An Overview*. MSFT Technical Report



# References

- A. Karatzoglou, X. Amatriain, L. Baltrunas, and N. Oliver. “*Multiverse recommendation: n-dimensional tensor factorization for context-aware collaborative filtering*”. In Proc. of the fourth ACM Recsys, 2010.
- S. Rendle, Z. Gantner, C. Freudenthaler, and L. Schmidt-Thieme. “*Fast context-aware recommendations with factorization machines*”. In Proc. of the 34th ACM SIGIR, 2011.
- S.H. Yang, B. Long, A.J. Smola, H. Zha, and Z. Zheng. “*Collaborative competitive filtering: learning recommender using context of user choice*”. In Proc. of the 34th ACM SIGIR, 2011.
- N. N. Liu, X. Meng, C. Liu, and Q. Yang. “*Wisdom of the better few: cold start recommendation via representative based rating elicitation*”. In Proc. of RecSys’11, 2011.
- M. Jamali and M. Ester. “*Trustwalker: a random walk model for combining trust-based and item-based recommendation*”. In Proc. of KDD ’09, 2009.

# References

- J. Noel, S. Sanner, K. Tran, P. Christen, L. Xie, E. V. Bonilla, E. Abbasnejad, and N. Della Penna. “*New objective functions for social collaborative filtering*”. In Proc. of WWW ’12, pages 859–868, 2012.
- X. Yang, H. Steck, Y. Guo, and Y. Liu. “*On top-k recommendation using social networks*”. In Proc. of RecSys’12, 2012.
- Dmitry Lagun. “*Modeling User Attention and Interaction on the Web*” 2014 PhD Thesis (Emory Un.)
- “*Robust Models of Mouse Movement on Dynamic Web Search Results Pages* - F. Diaz et al. In Proc. of CIKM 2013
- Amr Ahmed et al. “*Fair and balanced*”. In Proc. WSDM 2013
- Deepak Agarwal. 2013. *Recommending Items to Users: An Explore Exploit Perspective*. CIKM ‘13

# Online resources

- Recsys Wiki: <http://recsyswiki.com/>
- Recsys conference Webpage: <http://recsys.acm.org/>
- Recommender Systems Books Webpage: <http://www.recommenderbook.net/>
- Mahout Project: <http://mahout.apache.org/>
- MyMediaLite Project: <http://www.mymedialite.net/>



# Thanks!

## Questions?

Xavier Amatriain

xavier@netflix.com



@xamat



Xavier Amatriain – October 2014 – Recsys